Analysis and Design of Scalable Programming Auto grader

MTP Final Thesis Presentation

Guide:
Prof Varsha Apte

Submitted by:
Chaitanya Varma
203050026

Table of Contents

Need for Scalability in Programming auto-graders

Typical programming auto grader architecture

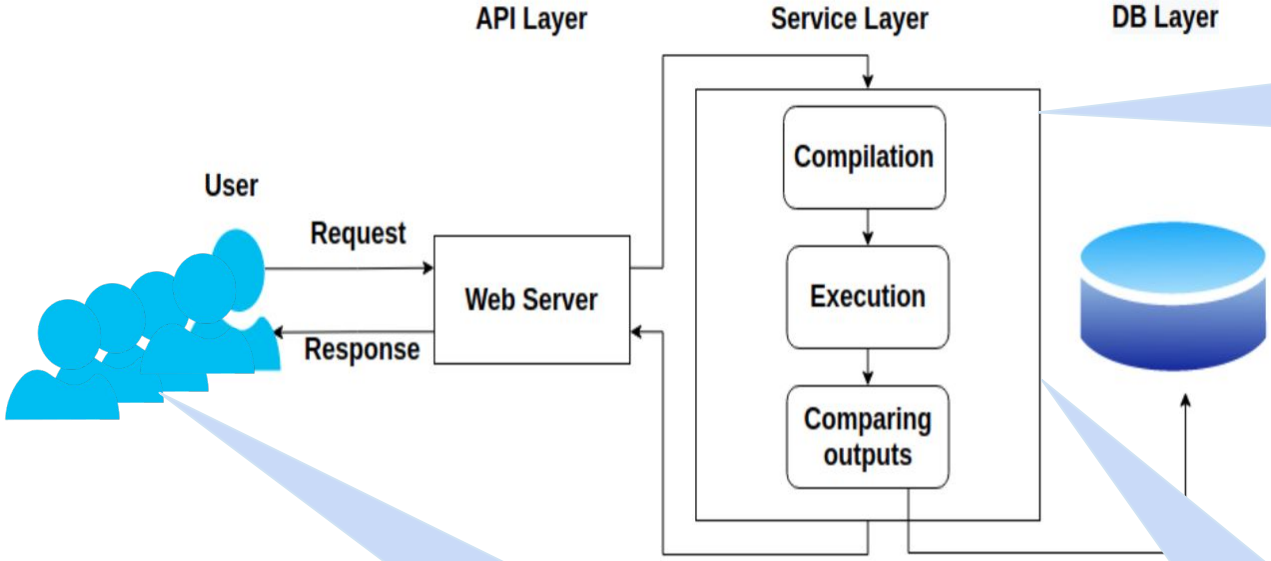

processing time depends upon the user code and cannot be profiled

Scalability: Ability of the system to scale its performance with increase in resources

At peak usage, system need to scale its performance

Not possible to estimate resource requirement in prior

Background about Evalpro

Evalpro Design

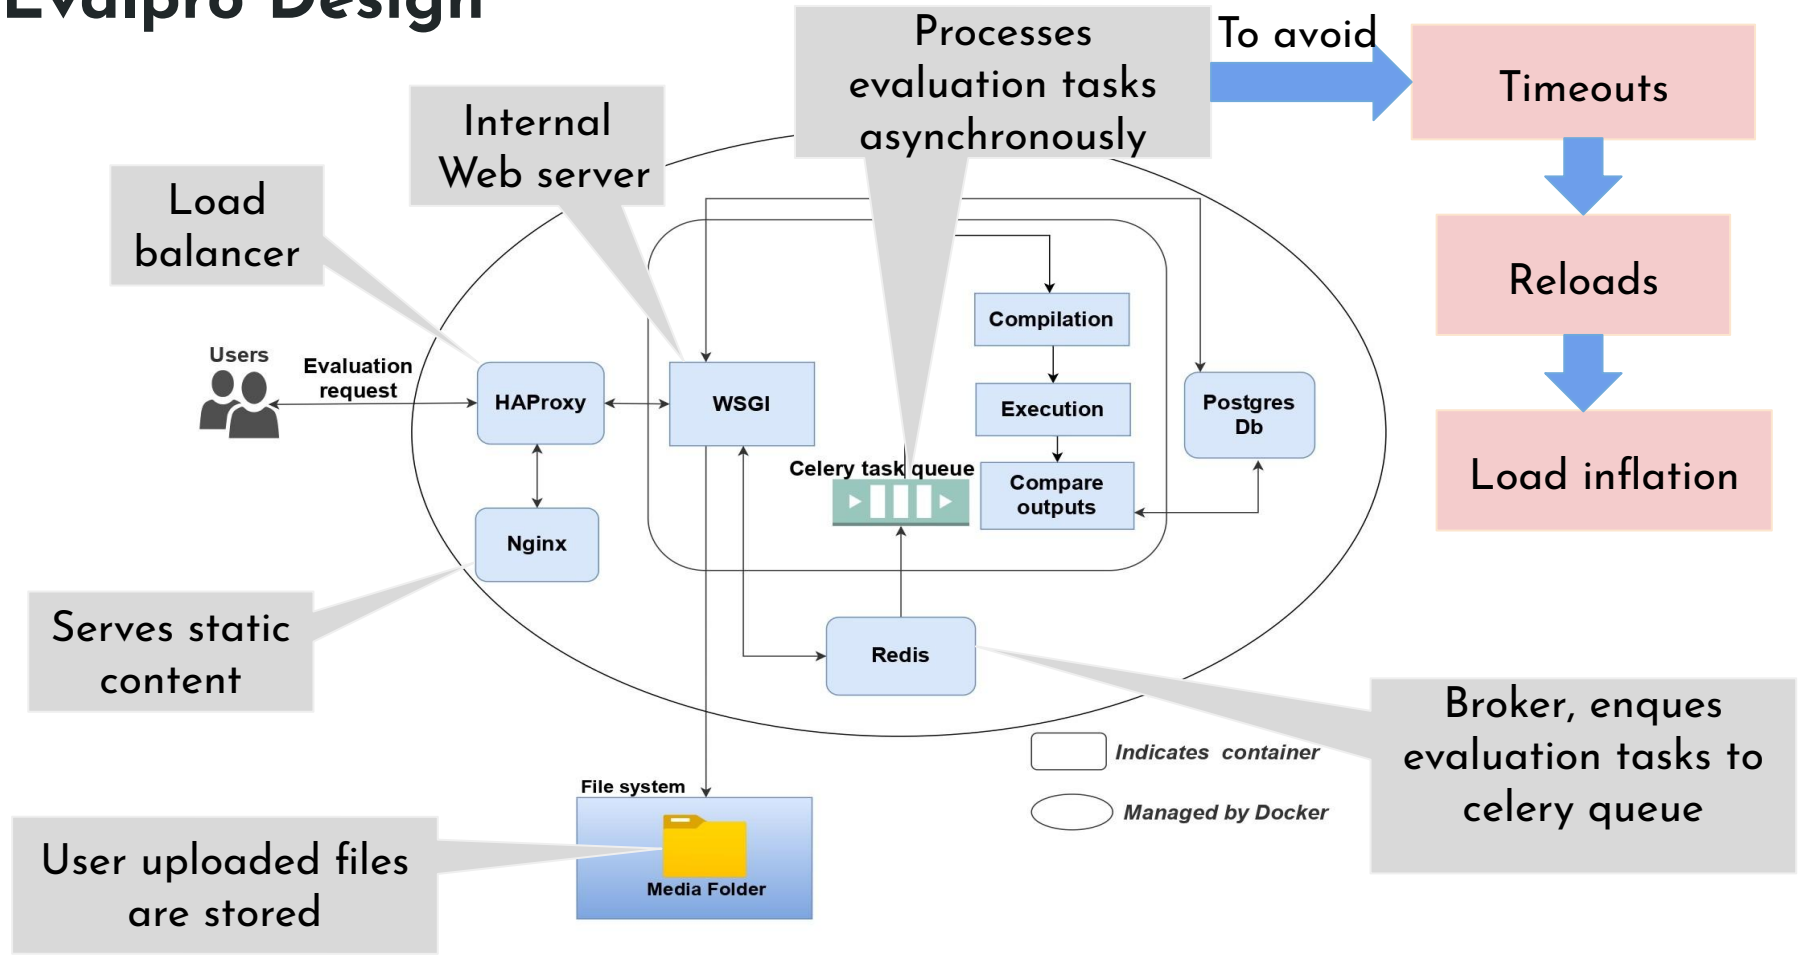

Prior work of Evalpro scalability

Prior work (MTP 2019) and its limitations

40 CPU core server

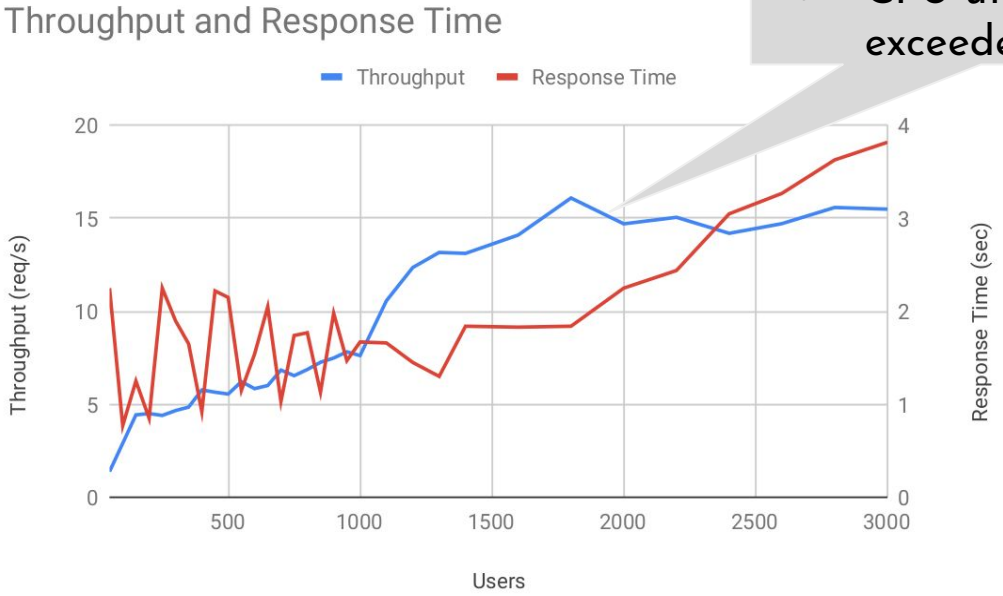

- Throughput flattened
- CPU utilization not exceeded 5%

Scalability issue

Increasing WSGI+Celery replicas increased the throughput and CPU utilization

- Reason for improvement in scalability not justified
- Scalability issue with single replica, configurations affecting the performance not studied

"The data and figures in this slide were taken from Anshul MTP thesis 2019"

Work Outline

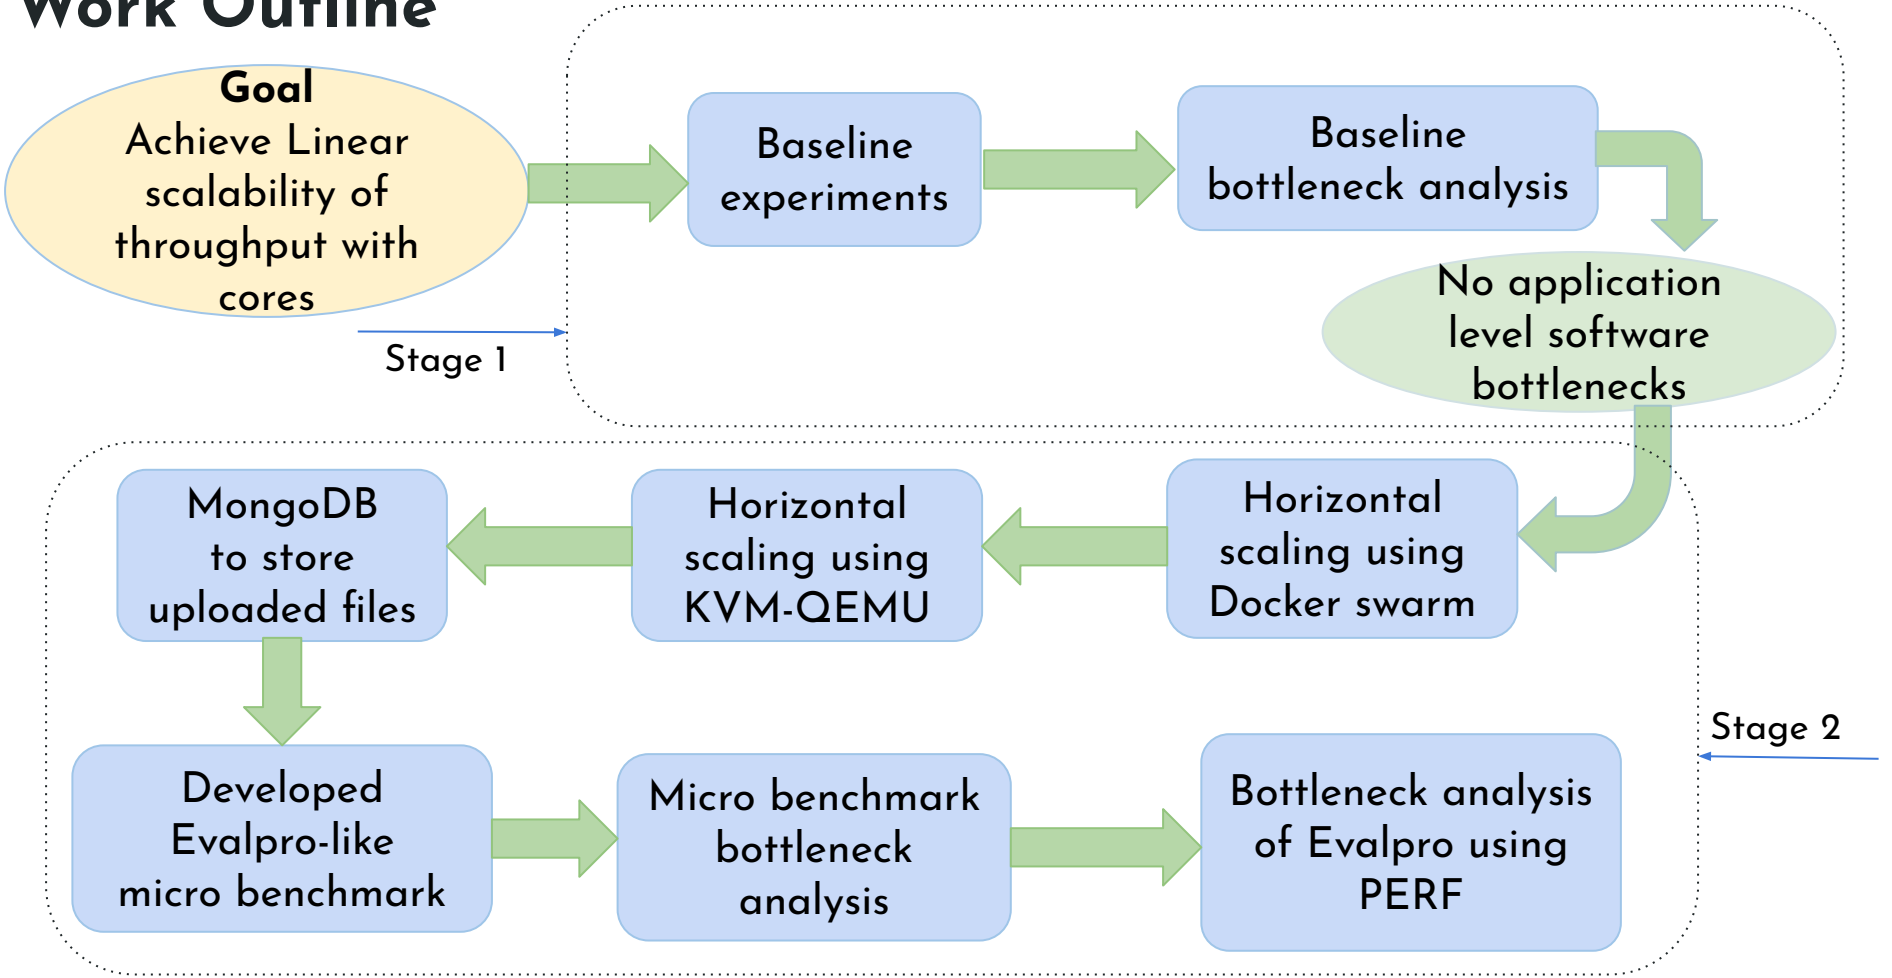

MTP stage 1

Load generation and performance measurement Infrastructure

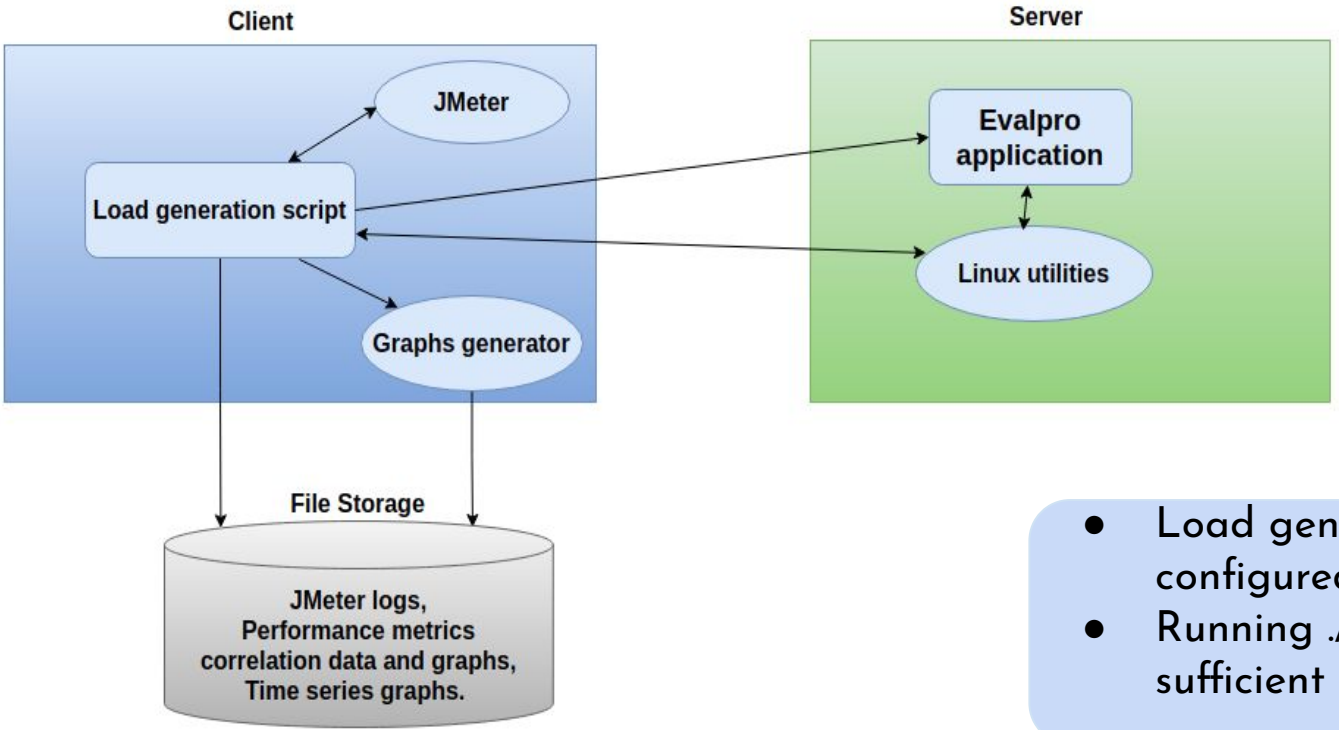

- Load generation configured for 5 minutes
- Running ./load_test.sh is sufficient

Baseline Experiment setup

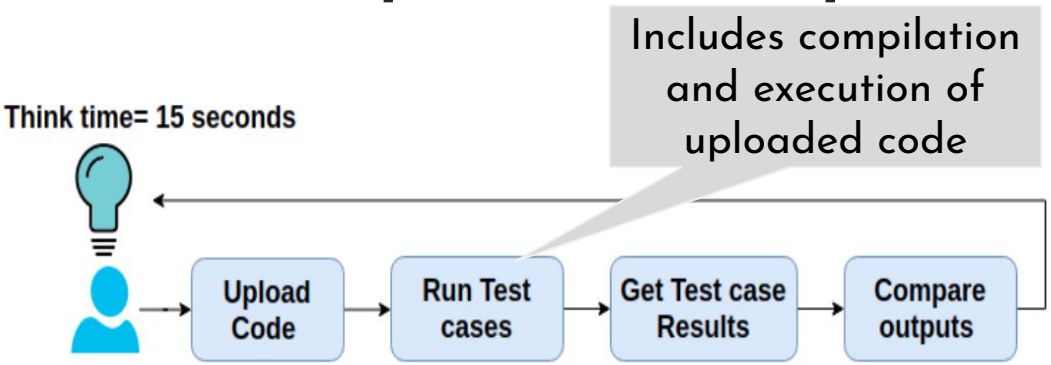

Realistic user session

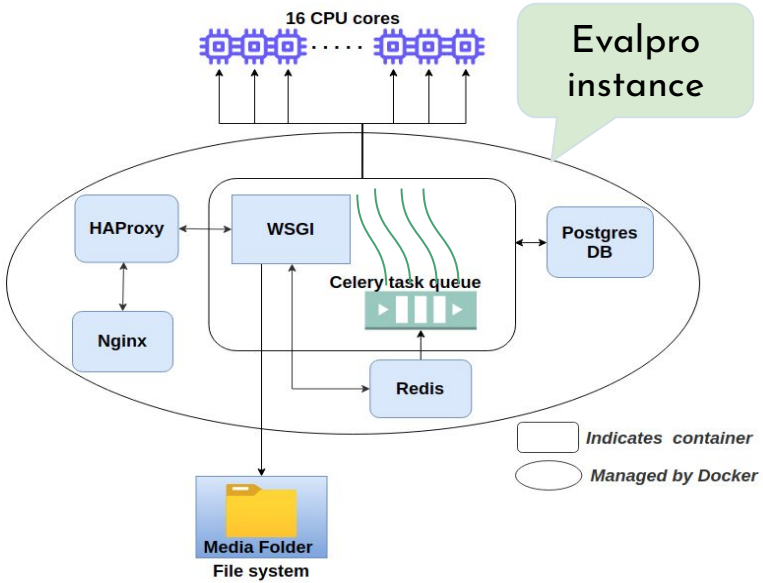

Type	CPU	Cores	Memory	L3 cache	L2 cache	L1 cache
Server	Intel^{R} Xeon^{R} CPU E5-2650 v2 @ 2.60GHz	16	16GB	20MB	256KB	32KB
Client	AMD Opteron^{TM} Processor 6212	16	16GB	6MB	2MB	64KB

Hardware specifications for baseline experiments

Baseline 16

To avoid obvious bottlenecks,
Postgres DB connections set to 10000
Celery threads set to number of cores

$Throughput_{max}(N)$: Maximum throughput with N CPU cores
$S(N)$: Throughput scalability factor with N CPU cores

$S(N) = Throughput_{max}(N) \div Throughput_{max}(1)$
Ideal $S(N) = N$

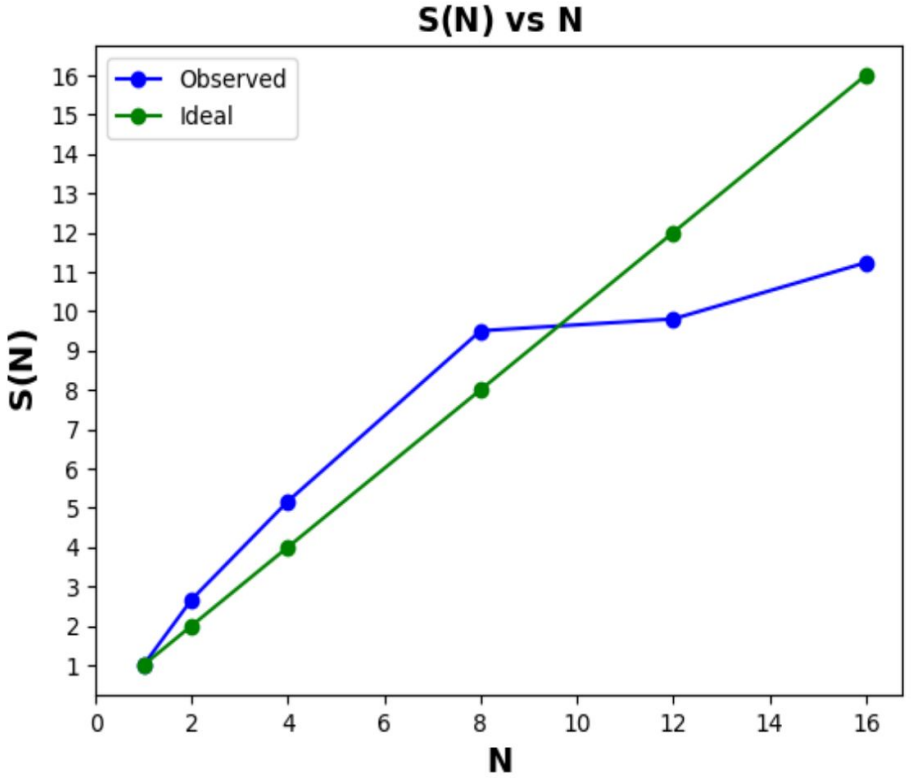

$Throughput_{max}(1)$	$Throughput_{max}(16)$	Ideal $Throughput_{max}(16)$	$S(16)$
0.3 req/sec	3.37 req/sec	4.8 req/sec	11.23

Linear scaling of throughput not achieved

Baseline bottleneck analysis

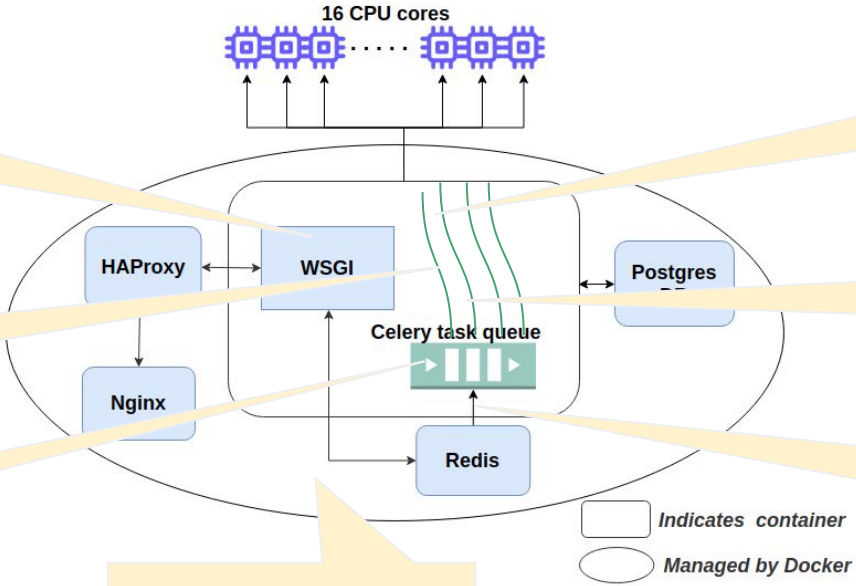

Increased gunicorn workers

Setting celery CPU affinities

Increased prefetch count

Increased celery threads

Made celery queue transient

Disabled logging

Increased open connections to redis

None of these experiments improved throughput scalability

No application level software bottlenecks

All cores 90-100 % utilized

MTP Stage-2

Horizontal scaling with Docker swarm

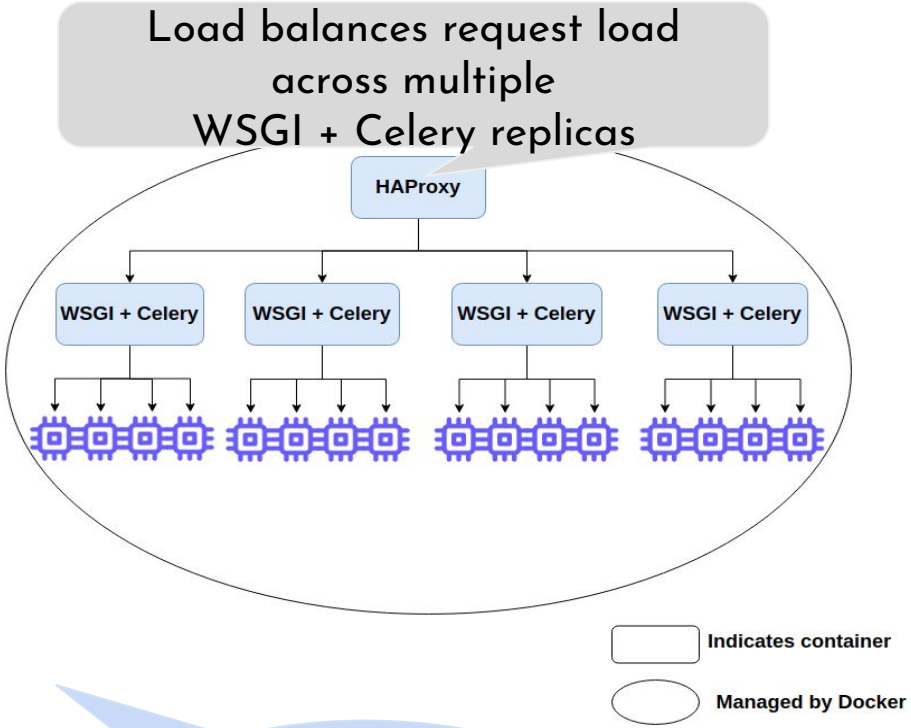

	Baseline	Using Docker swarm
$Throughput_{max}(1)$ (req/sec)	0.3	0.3
$Throughput_{max}(16)$ (req/sec)	3.37	3.4
Ideal $Throughput_{max}(16)$ (req/sec)	4.8	4.8
$S(16)$	11.23	11.3

Docker swarm didn't improve throughput scalability

Horizontal scaling with KVM-QEMU

Completely Isolated VM setup

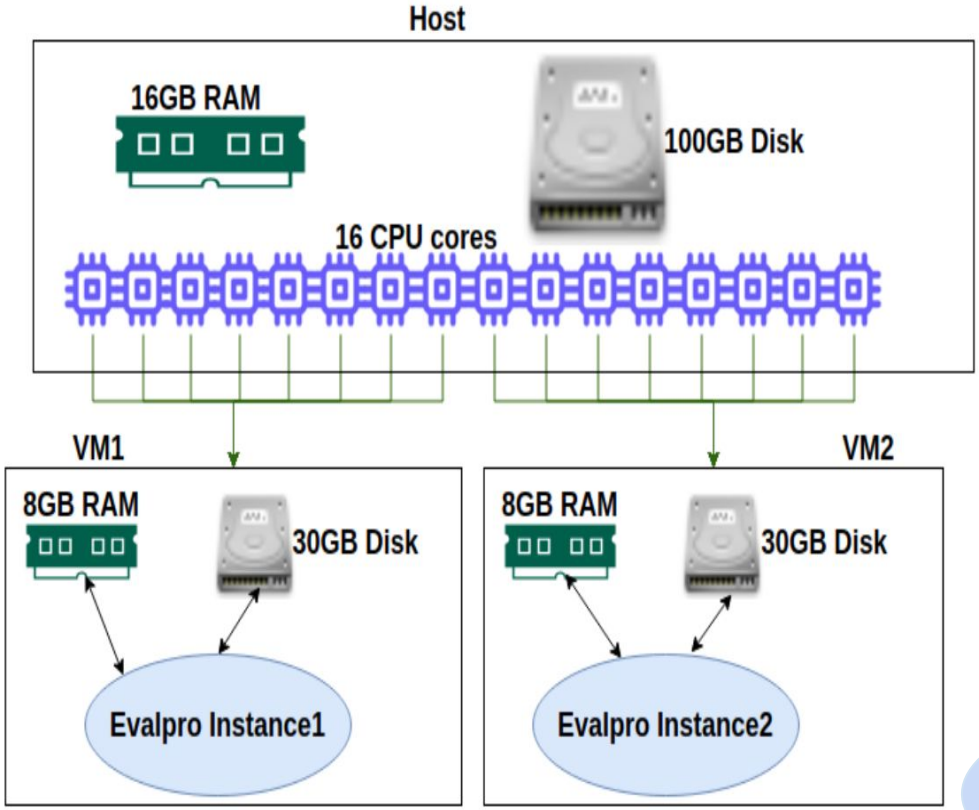

	Baseline	Isolated VM setup
$Throughput_{max}(1)$ (req/sec)	0.3	0.3
$Throughput_{max}(16)$ (req/sec)	3.37	4.45
Ideal $Throughput_{max}(16)$ (req/sec)	4.8	4.8
$S(16)$	11.23	14.5

Isolated setup improved throughput scalability

Completely Isolated VM setup limitations

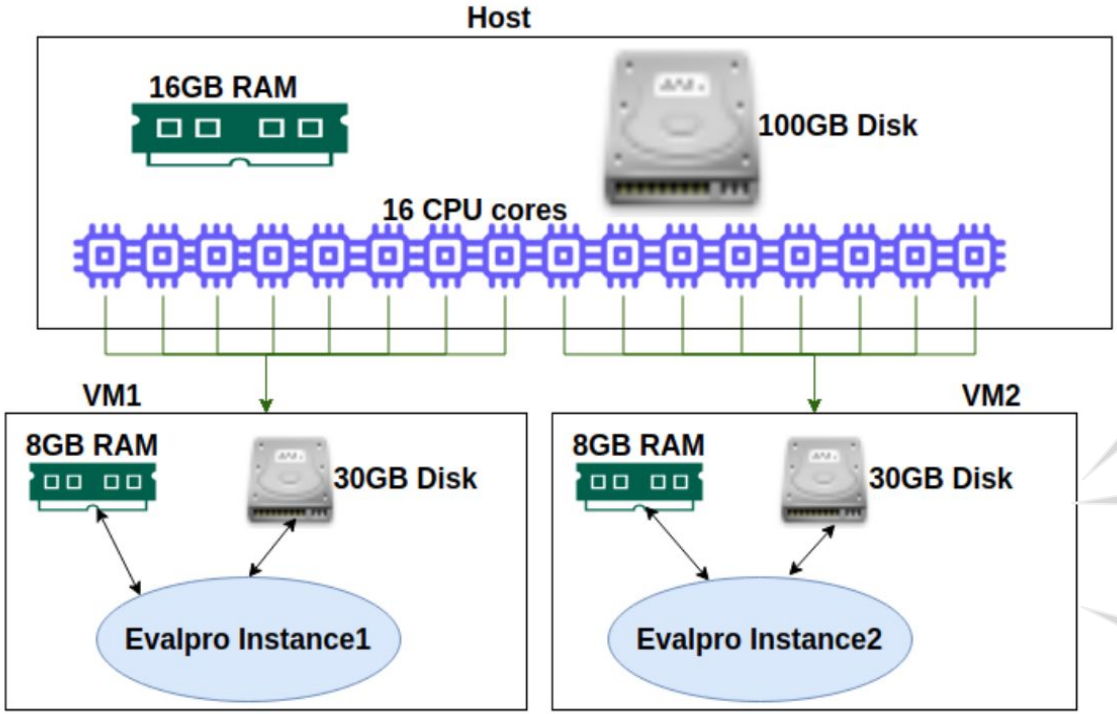

Each VM need to maintain separate user related data

Load balancer need have user to VM mapping

Migrating user data with new VM addition is very complex and not feasible.

Not the desired solution

User data and files sharing VM setup

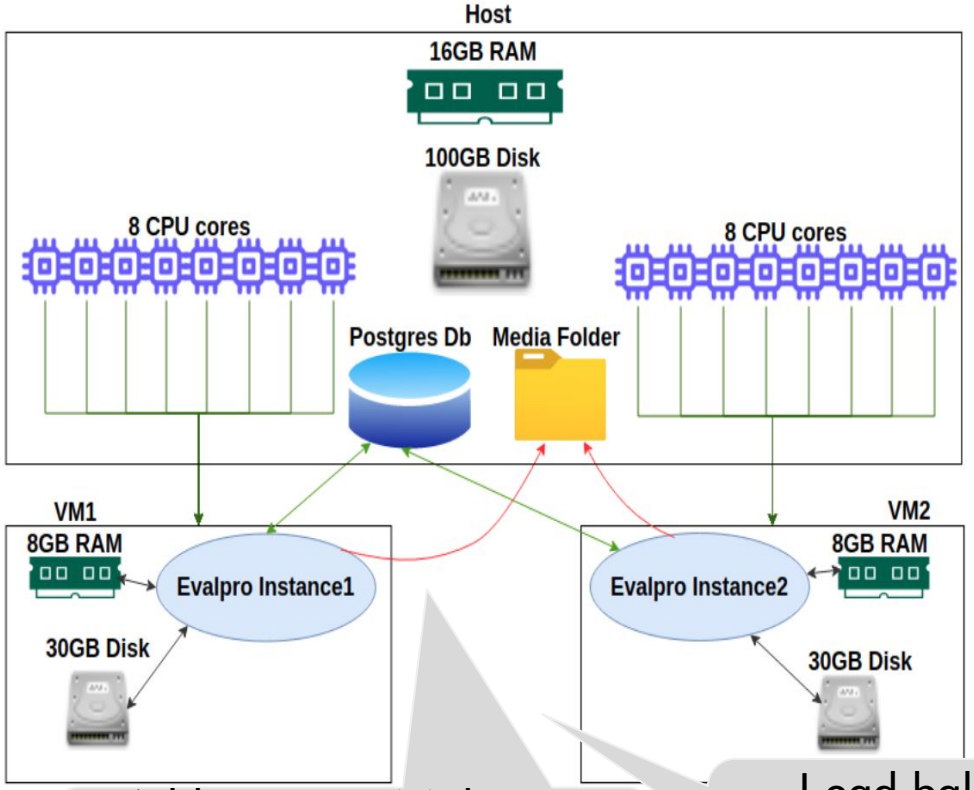

	Baseline	Data and file sharing VM setup
$Throughput_{max}(1)$ (req/sec)	0.3	0.3
$Throughput_{max}(16)$ (req/sec)	3.37	2.78
Ideal $Throughput_{max}(16)$ (req/sec)	4.8	4.8
$S(16)$	11.23	9.62

Adding new VM doesn't require migration of user data

Load balancer need not maintain user to VM mapping

Throughput scalability decreased with this setup

Reason for reduction in throughput scalability

Performed experiments with different user sessions

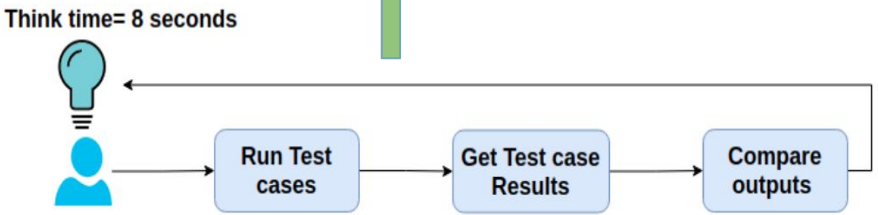

User session: Without Upload

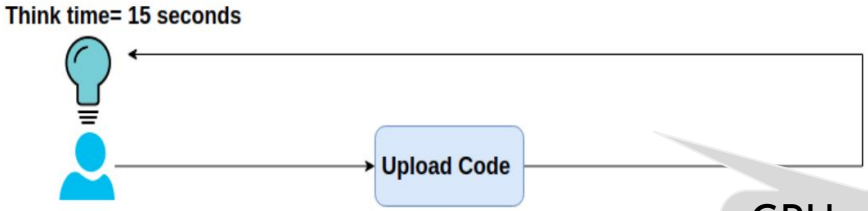

User session: Only Upload

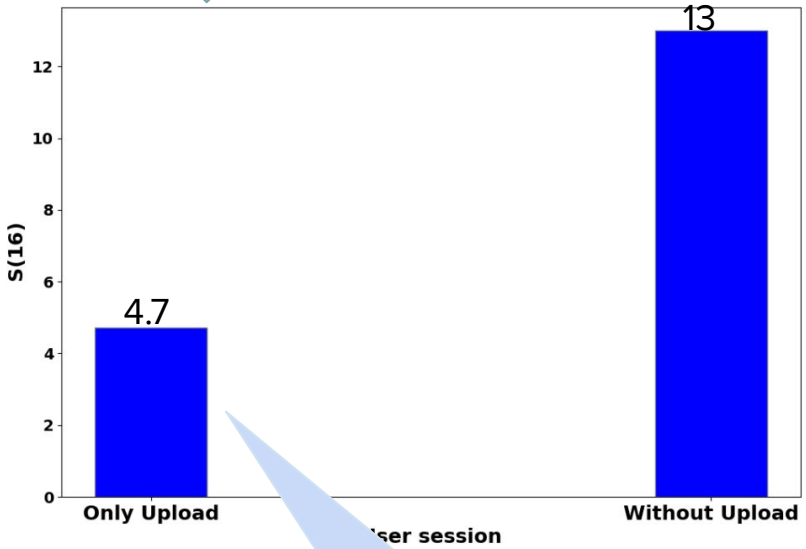

CPU utilization never exceeded 35%

Uploading file by VMs in host is reducing scalability

Using MongoDB for File Storage

Evalpro updated architecture

Stores user uploaded files

File data is stored as stream of bytes

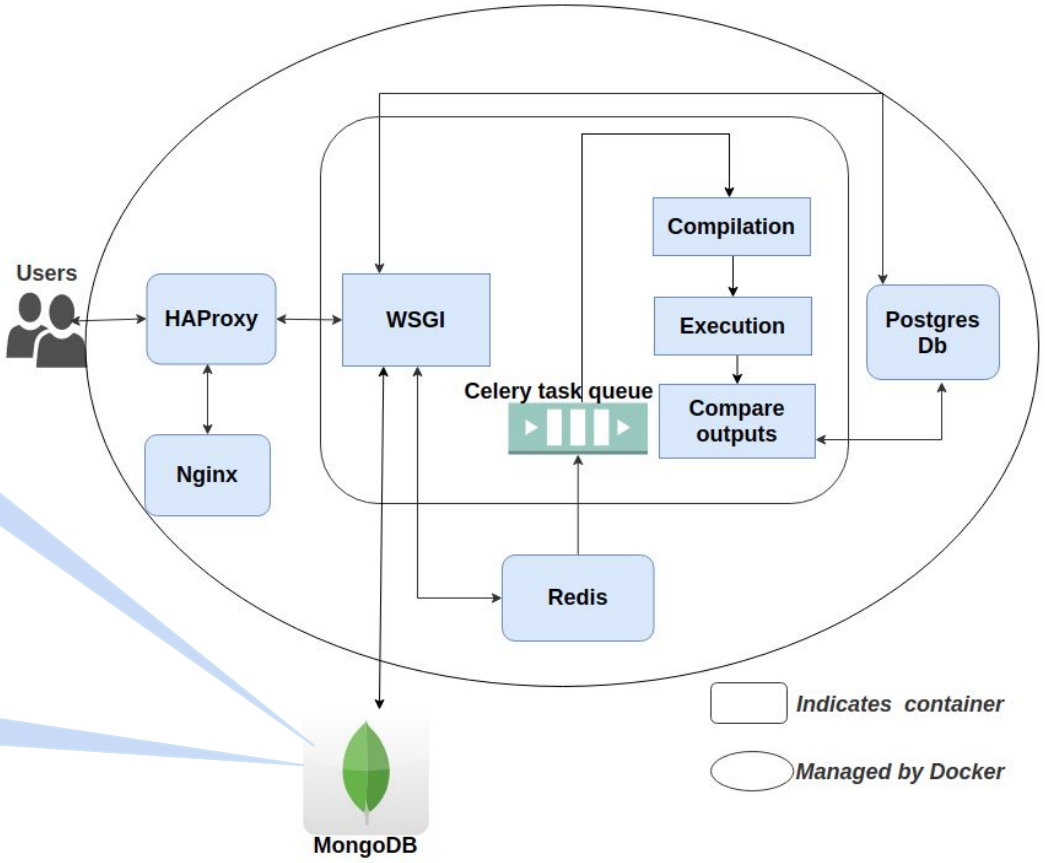

User data and files sharing VM setup with MongoDB

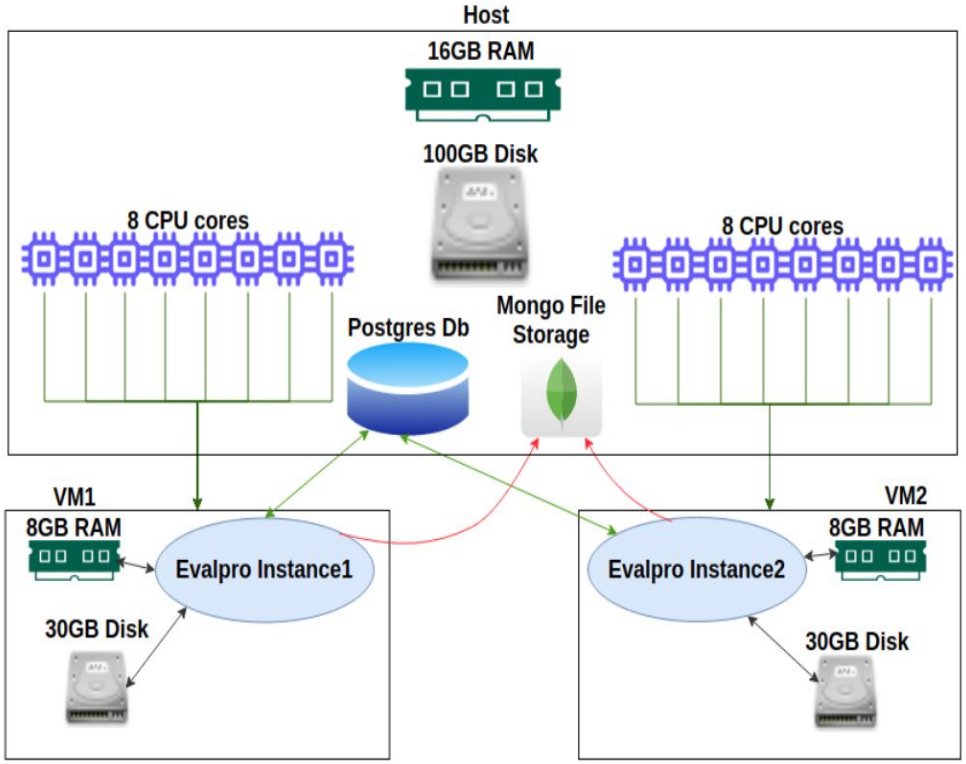

	Baseline	MongoDB for file storage
$Throughput_{max}(1)$ (req/sec)	0.3	0.29
$Throughput_{max}(16)$ (req/sec)	3.37	3.5
Ideal $Throughput_{max}(16)$ (req/sec)	4.8	4.64
$S(16)$	11.23	12

Hypothesis:
With large number of CPU cores throughput scalability will further increase

Throughput scalability slightly increased

Experiment setup using 64 CPU cores

Type	CPU	Cores	Memory	L3 Cache	L2 cache	L1 cache
Server	Intel^{R} Xeon^{R} CPU E5-2683 v4 @ 2.10GHz	64	128GB	80MB	8MB	1MB
Client	AMD Opteron^{TM} Processor 6212	16	16GB	6MB	2MB	64KB
Client	AMD Opteron^{TM} Processor 6278	16	16GB	6MB	2MB	64KB
Client	Intel^{R} Xeon^{R} CPU E5-2650 v2 @ 2.60GHz	16	16GB	20MB	256KB	32KB

Baseline-64

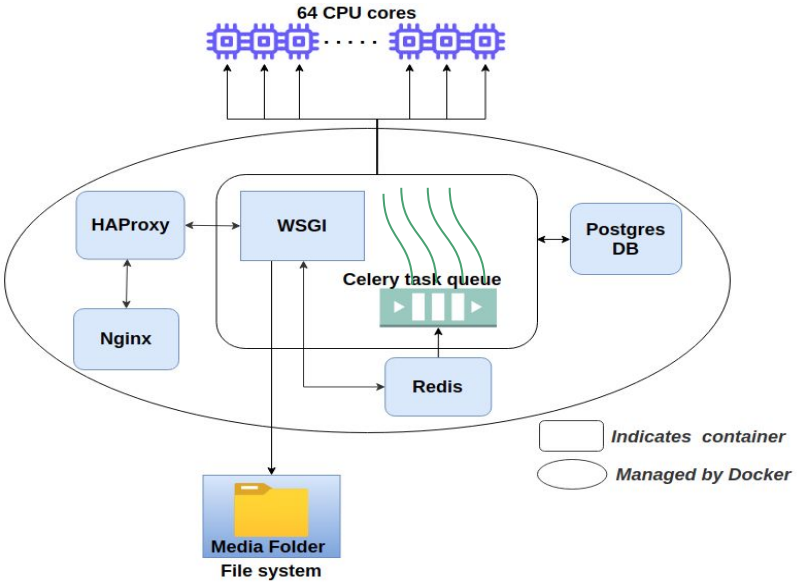

$Throughput_{max}(1)$	0.66 req/sec
$Throughput_{max}(16)$	8.75 req/sec
$Throughput_{max}(64)$	19.56 req/sec
Ideal $Throughput_{max}(64)$	42 req/sec
$S(64)$	29.6

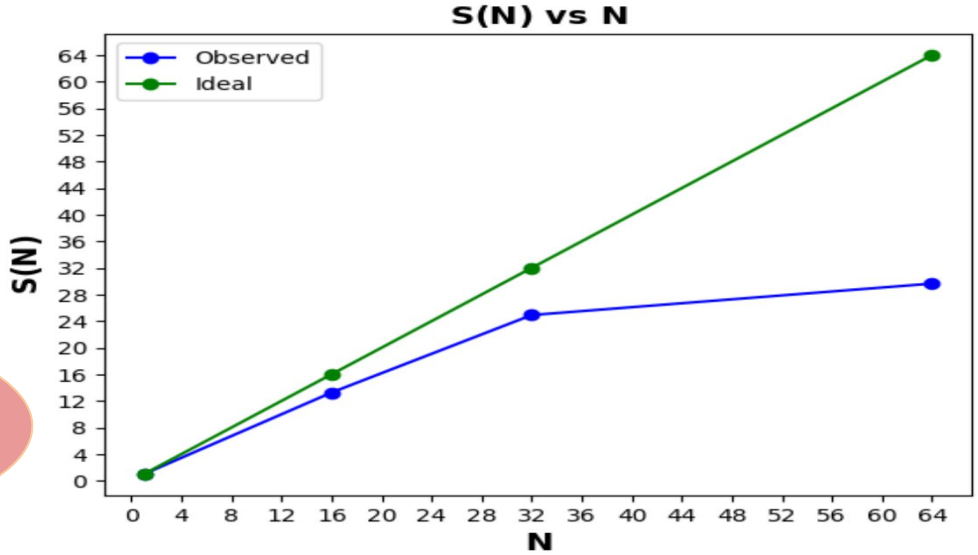

All CPU cores 90-100% utilized

Linear scaling of throughput not achieved

User data and files sharing VM setup with MongoDB - 64

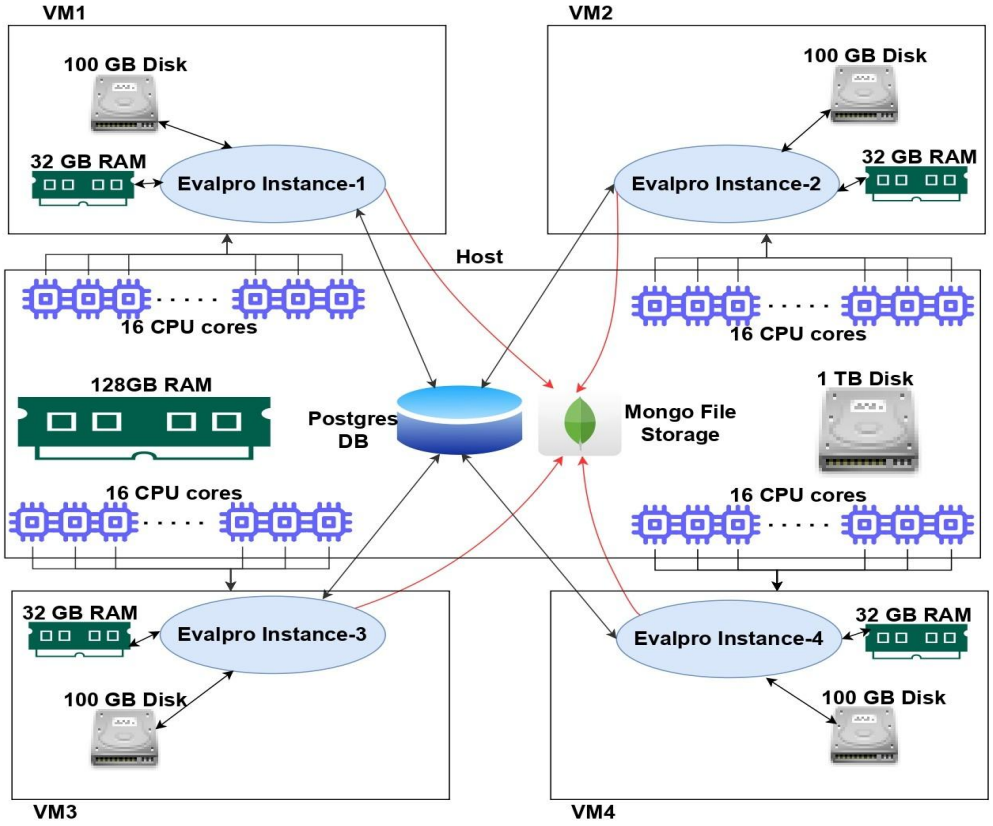

	Baseline-64	MongoDB for file storage - 64
$Throughput_{max}(1)$ (req/sec)	0.66	0.47
$Throughput_{max}(64)$ (req/sec)	19.56	11.27
Ideal $Throughput_{max}(64)$ (req/sec)	42	30
$S(64)$	29.6	24

Our hypothesis is wrong

Throughput scalability decreased

Completely Isolated VM setup - 64

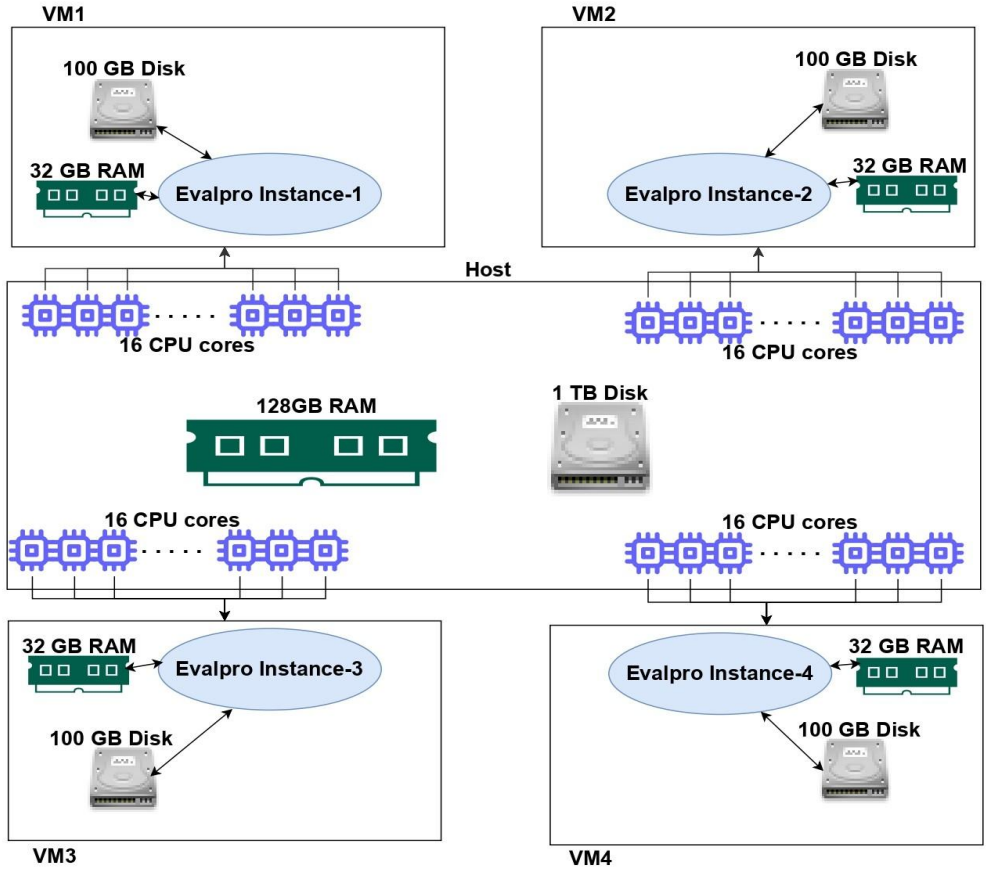

	Baseline-64	Isolated setup- 64
$Throughput_{max}(1)$ (req/sec)	0.66	0.66
$Throughput_{max}(64)$ (req/sec)	19.56	15.03
Ideal $Throughput_{max}(64)$ (req/sec)	42	42
	29.6	23

With higher number of cores, even this setup didn't improve throughput scalability

Throughput scalability decreased

Micro benchmark Experiments

Need for micro benchmarks

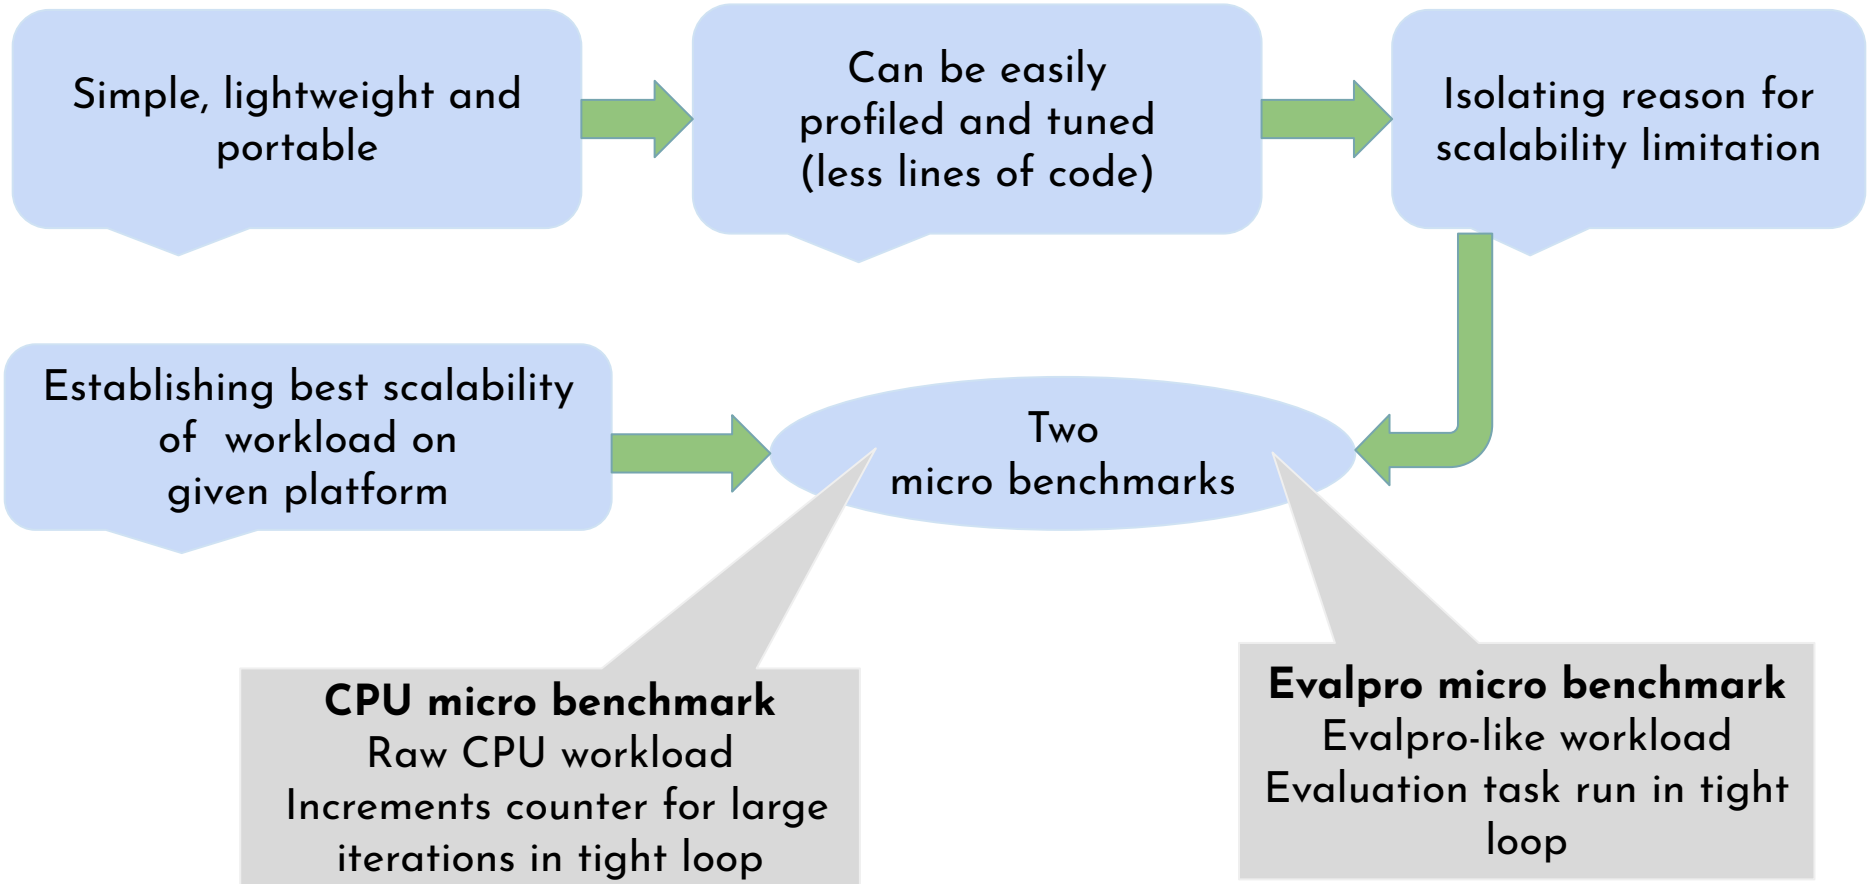

Experiments on CPU micro benchmark

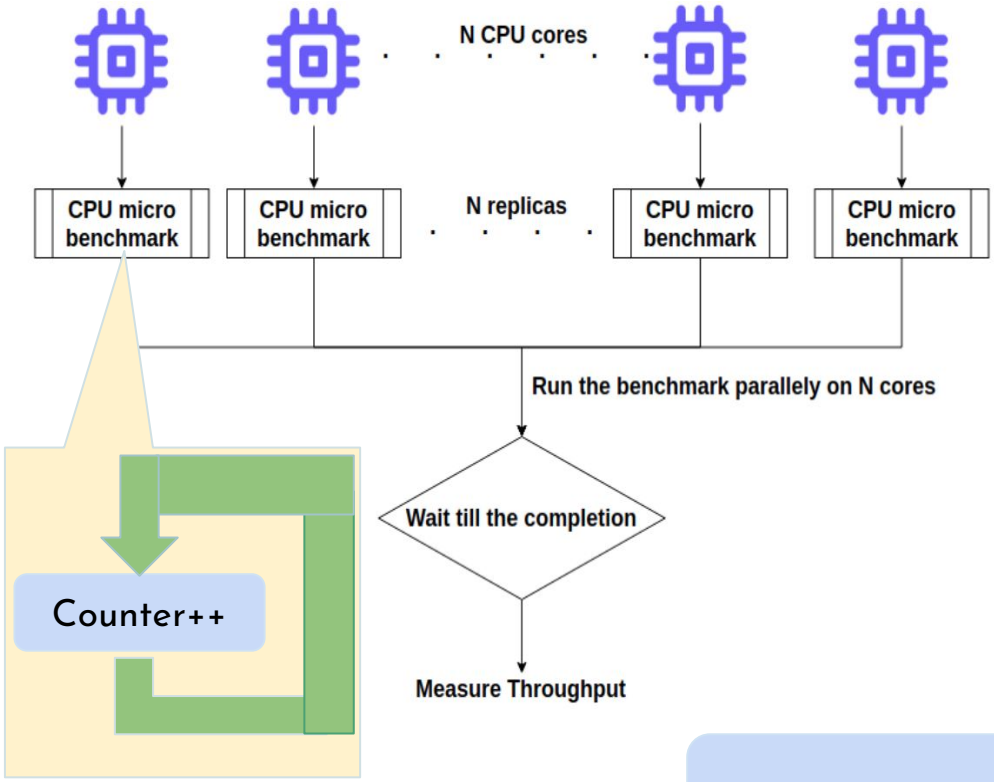

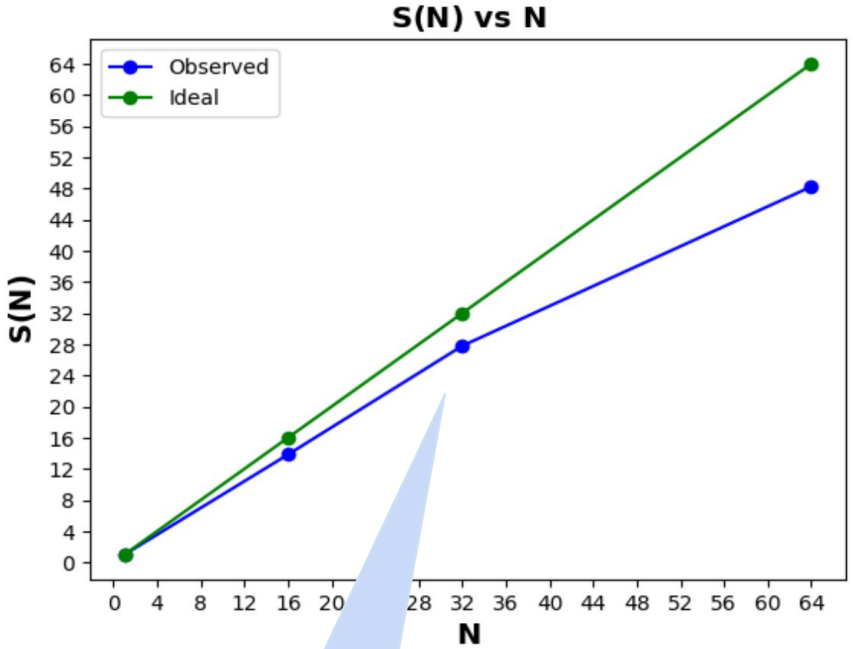

S(64) = 48

Throughput scalability is close to linear

Experiments on Evalpro micro benchmark

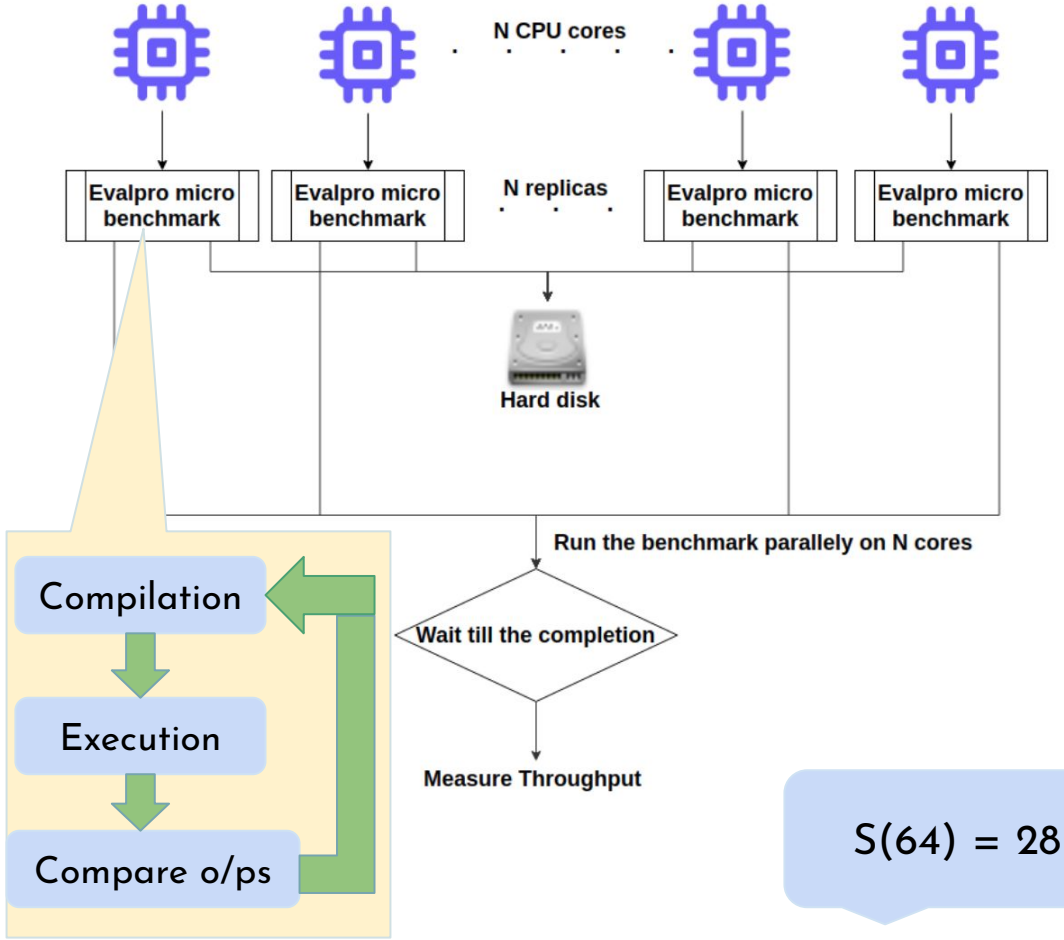

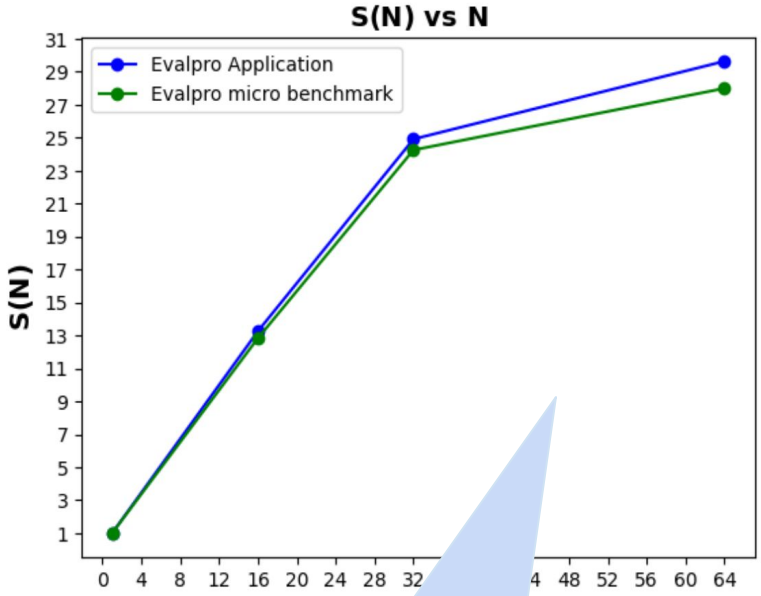

S(64) = 28

Throughput scalability is almost same as Evalpro application

Micro benchmark Bottleneck Analysis

Using Tmpfs in place of Hard disk

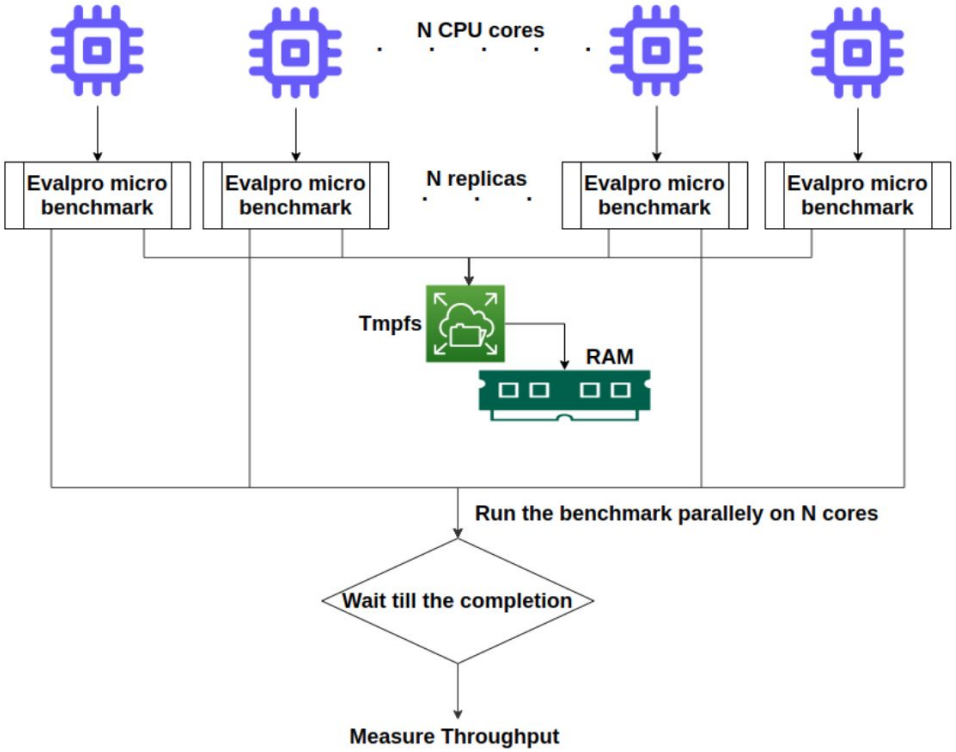

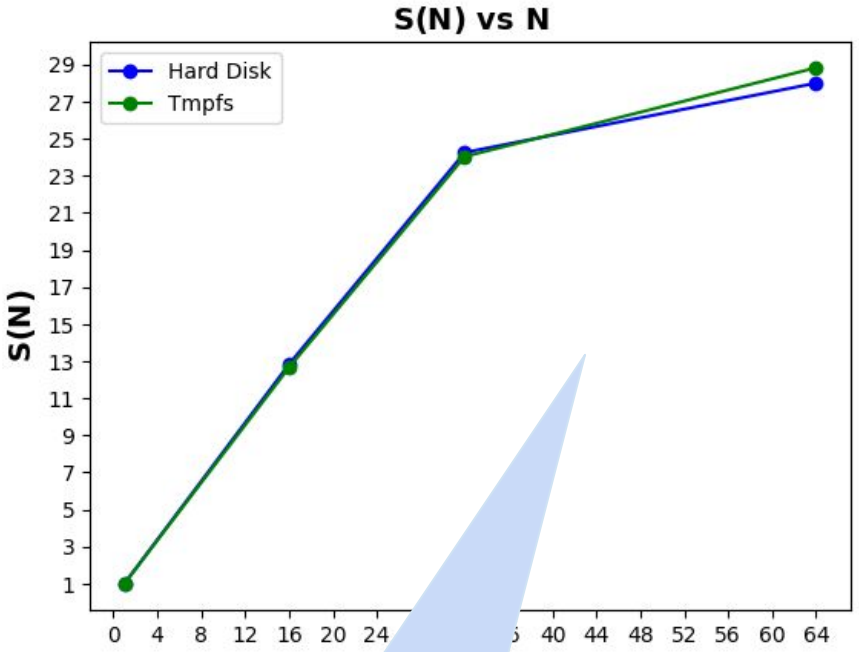

Using Tmpfs
didn't improve
throughput scalability

PERF analysis on Evalpro micro benchmark

PERF profiles
different software, hardware events

$Count_{event}(N)$: Occurrence count of event with N cores
$Inflation\_factor_{event}(N)$: Inflation factor of event with N cores

$$Inflation\_factor_{event}(N) = Count_{event}(N) \div (N \times Count_{event}(1))$$

Ideal $Inflation\_factor_{event}(N) = 1$

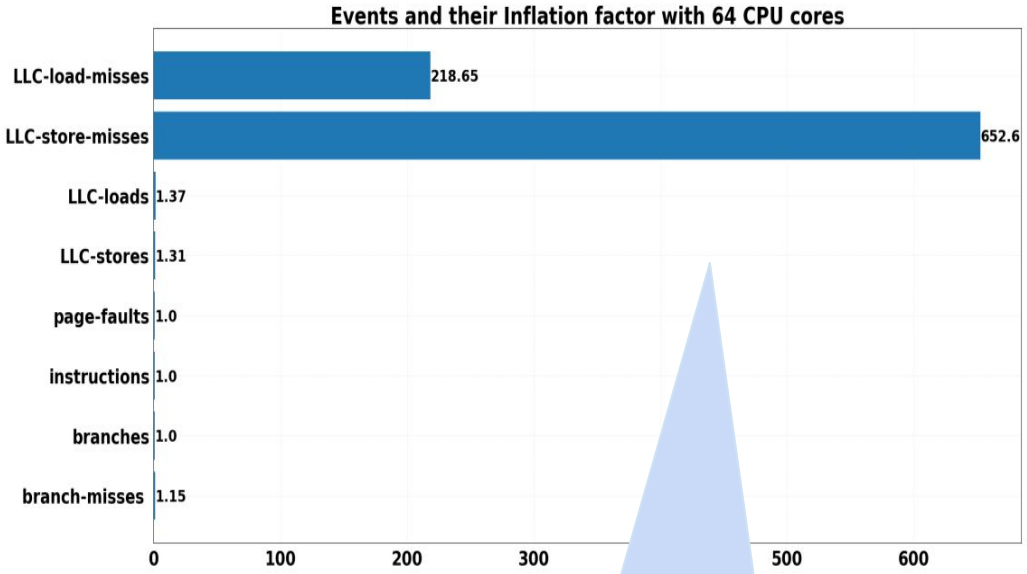

LLC load and store misses, disproportionately inflated with 64 CPU cores

PERF analysis at Program level

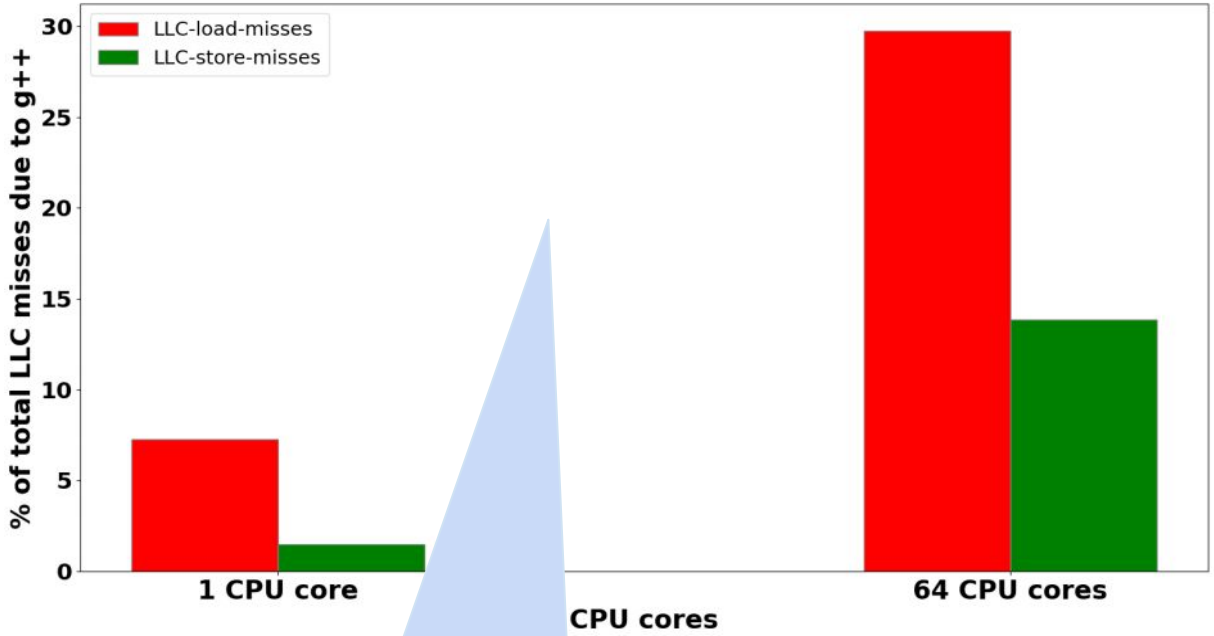

LLC load and store misses increased drastically by g++ when number of cores is high

Hypothesis
CPU cache is reason for scalability limitation of Evalpro application

CPU cache is the reason for scalability limitation of Evalpro micro benchmark

PERF analysis on Evalpro application

16 core

Baseline -16

20 MB CPU cache Intel^{R} Xeon^{R} CPU E5-2650 v2 @ 2.60GHz

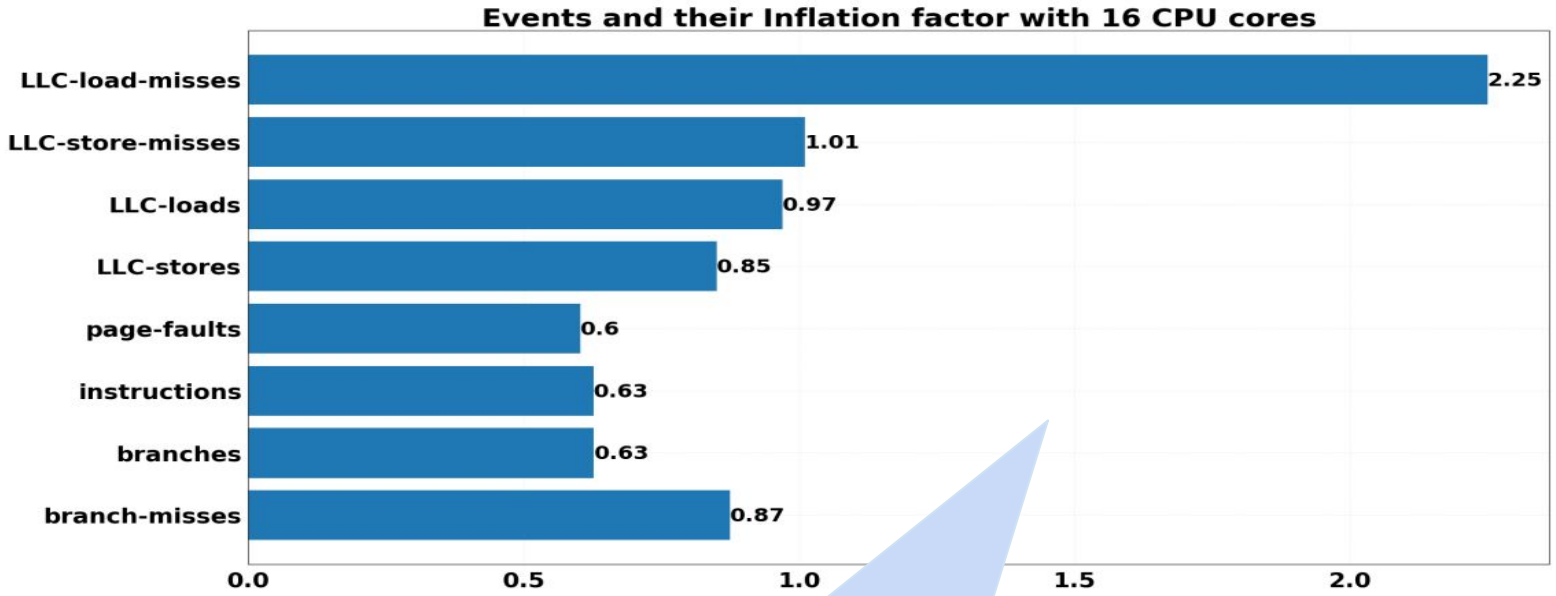

LLC load misses inflated
LLC store misses slightly inflated

Cache size affect on Evalpro application

Baseline - 16
16 cores
20 MB cache

Baseline -64
only 16 cores enabled
89 MB cache

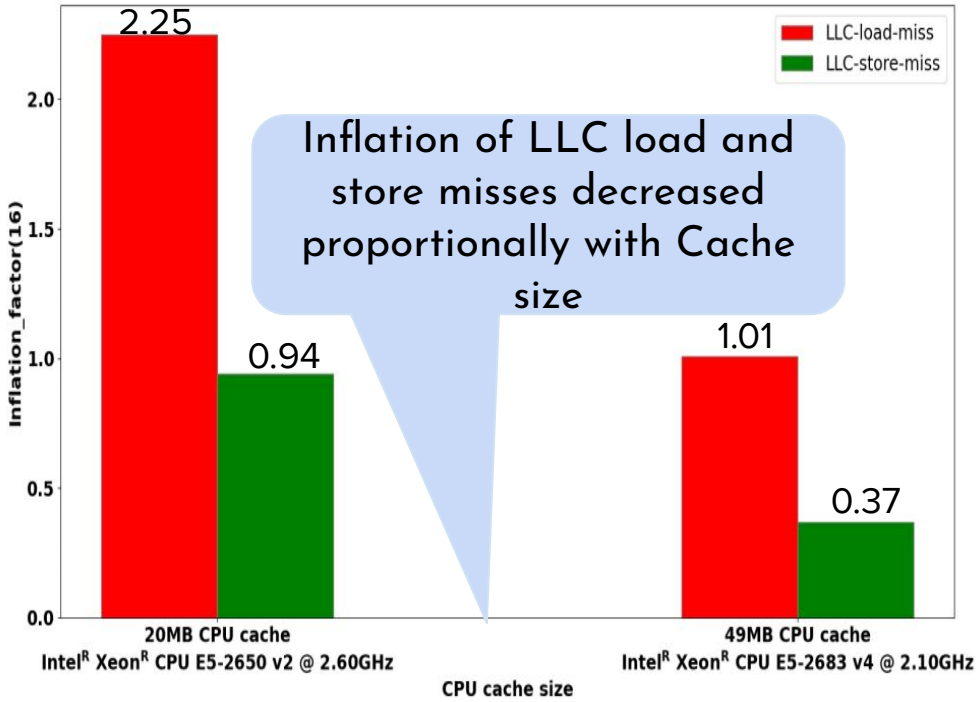

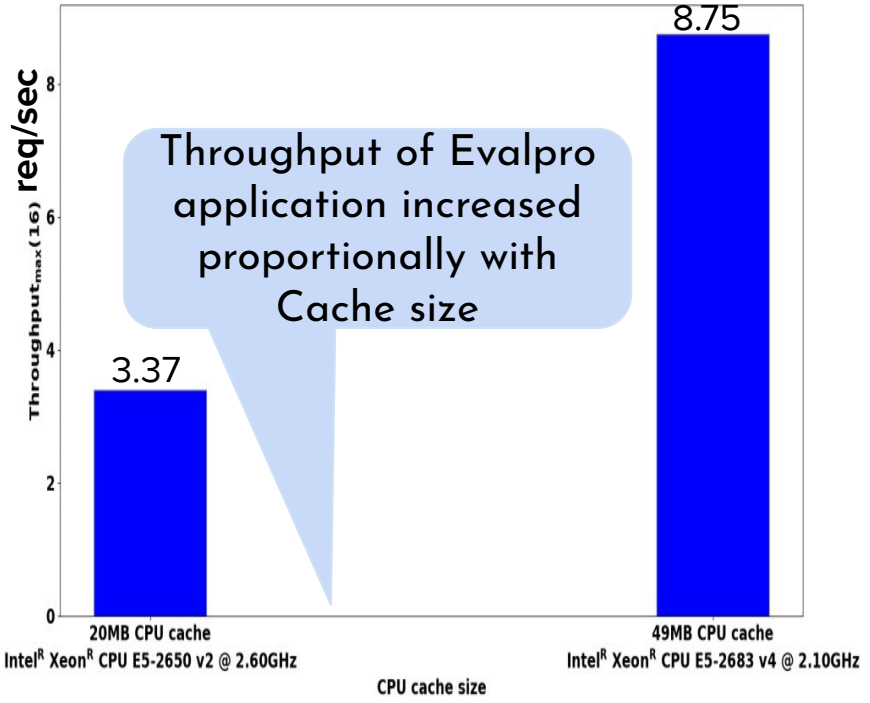

Cache size affect summary

Our hypothesis
is correct
i.e CPU Cache
Is the bottleneck

	Baseline-16	Baseline-64	% of increase or decrease
CPU cache size for 16 CPU cores	20 MB	49 MB	+ 145%
$Inflation\_factor_{LLC\text{-}load\text{-}misses}(16)$	2.25	0.94	- 140%
$Inflation\_factor_{LLC\text{-}store\text{-}misses}(16)$	1.01	0.37	- 170%
$Throughput_{max}(16)$	3.37 req/sec	8.75 req/sec	+ 160%

Conclusion

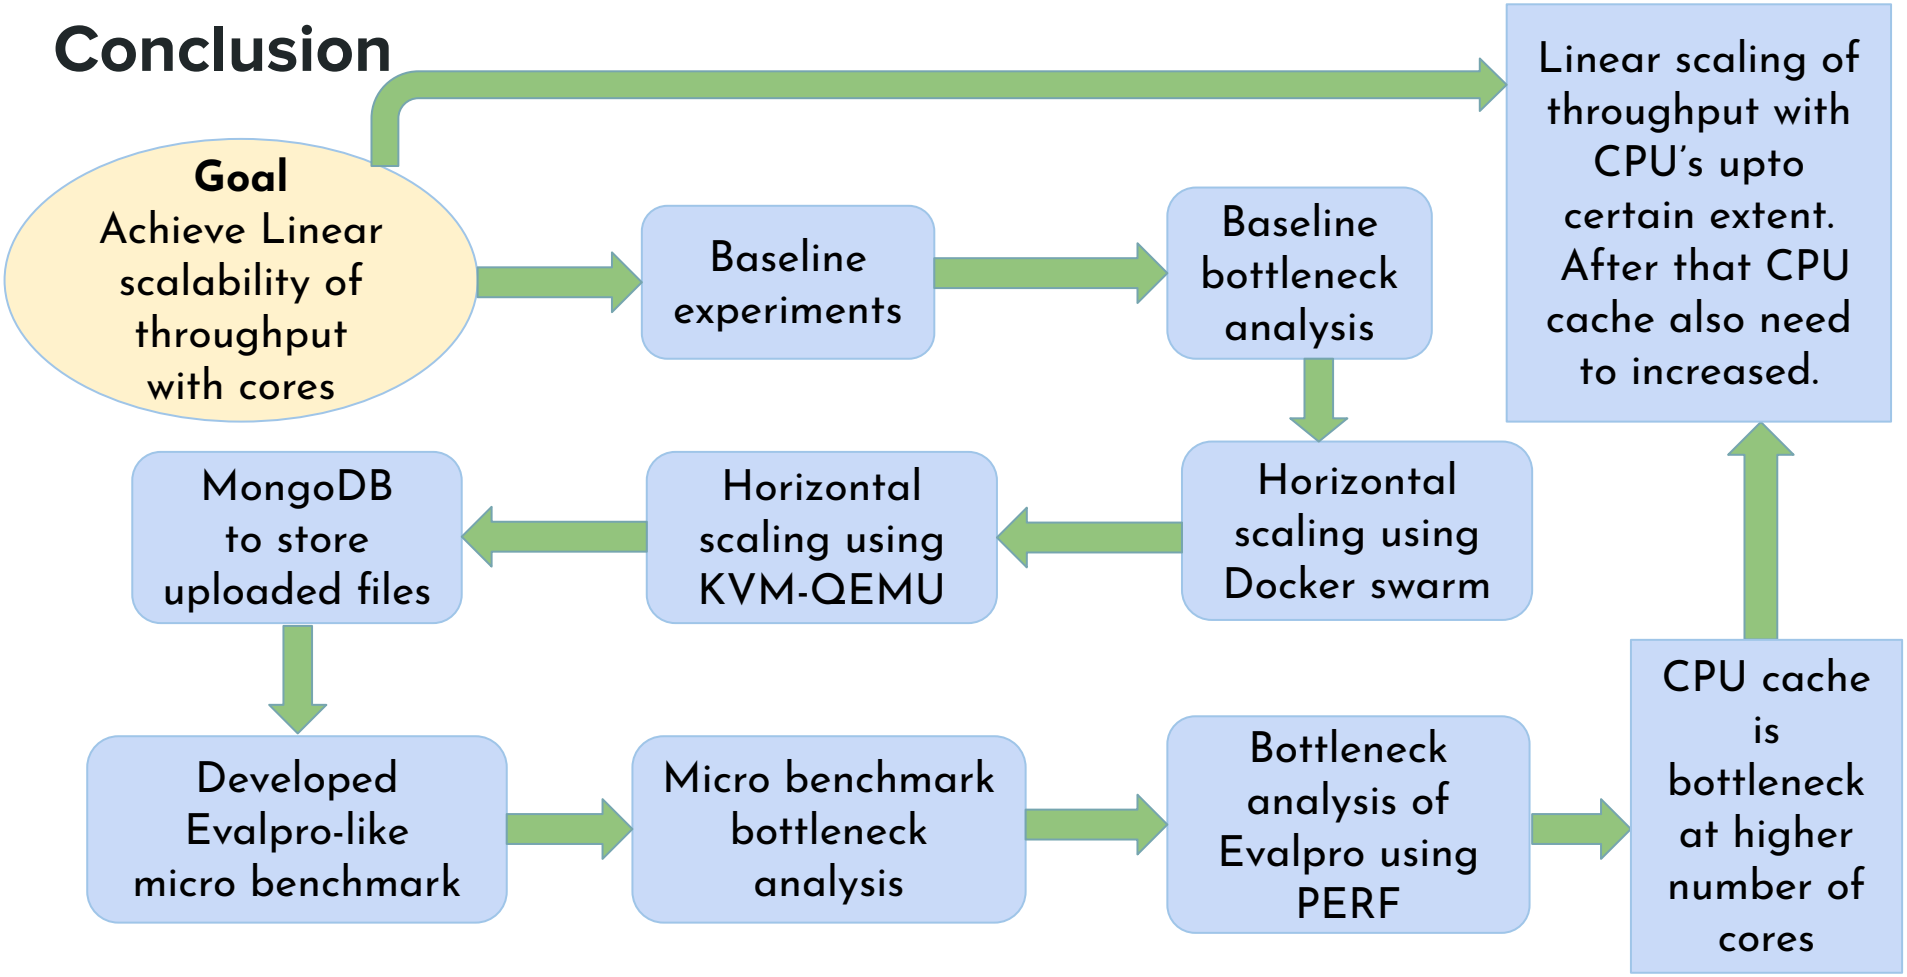
Goal
Achieve Linear scalability of throughput with cores
Baseline experiments
Baseline bottleneck analysis
Linear scaling of throughput with CPU's upto certain extent. After that CPU cache also need to increased.
Horizontal scaling using Docker swarm
Horizontal scaling using KVM-QEMU
MongoDB to store uploaded files
Developed Evalpro-like micro benchmark
Micro benchmark bottleneck analysis
Bottleneck analysis of Evalpro using PERF
CPU cache is bottleneck at higher number of cores

Final Recommendation for Scaling EvalPro

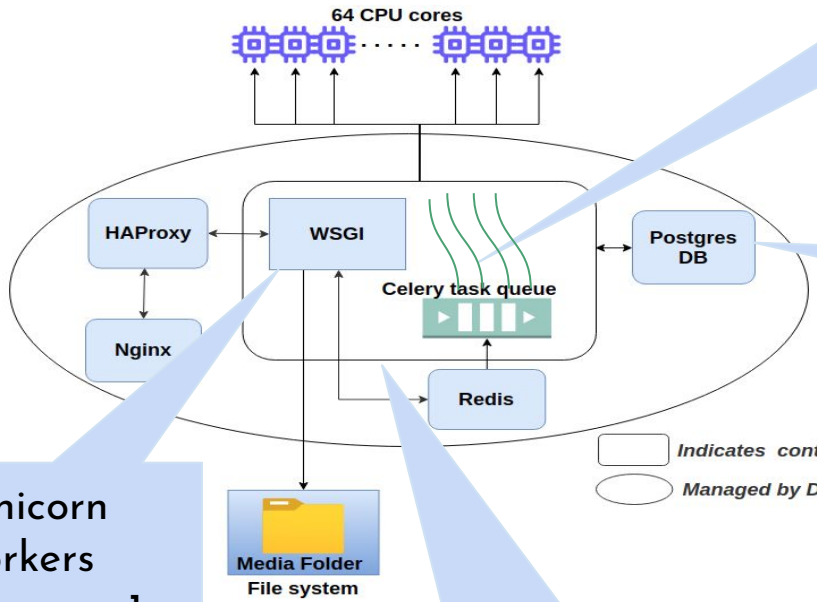

celery threads at least number of cores

postgres connections to large number i.e around 100000

Gunicorn workers = 2*ncores+1

No need for multiple WSGI+Celery replicas with above configurations

Migrate Evalpro to cloud

Pay for use

Easily scale performance linearly with CPU cores and cache size

Future work

Scalability of compilation workloads

Developing a browser plugin to do compilation at client side

Improving single core throughput of Evalpro

THANK YOU

Backup Slides

Increasing gunicorn workers

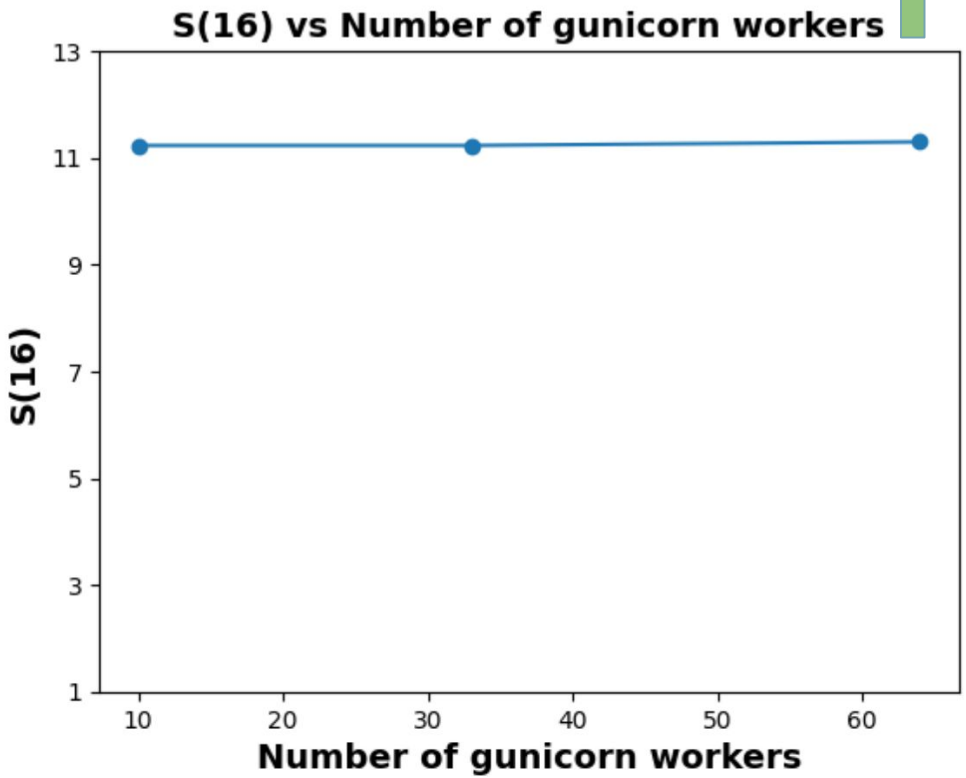

Increasing gunicorn workers didn't improve throughput scalability

Gunicorn workers are WSGI worker threads

Baseline gunicorn workers = 10

Gunicorn workers are not bottleneck

Gunicorn documentation says,
Ideal value = 2 x ncores +1

Tuning celery threads

Auto grading tasks processed asynchronously

Celery threads are bottleneck when less than number of CPU cores

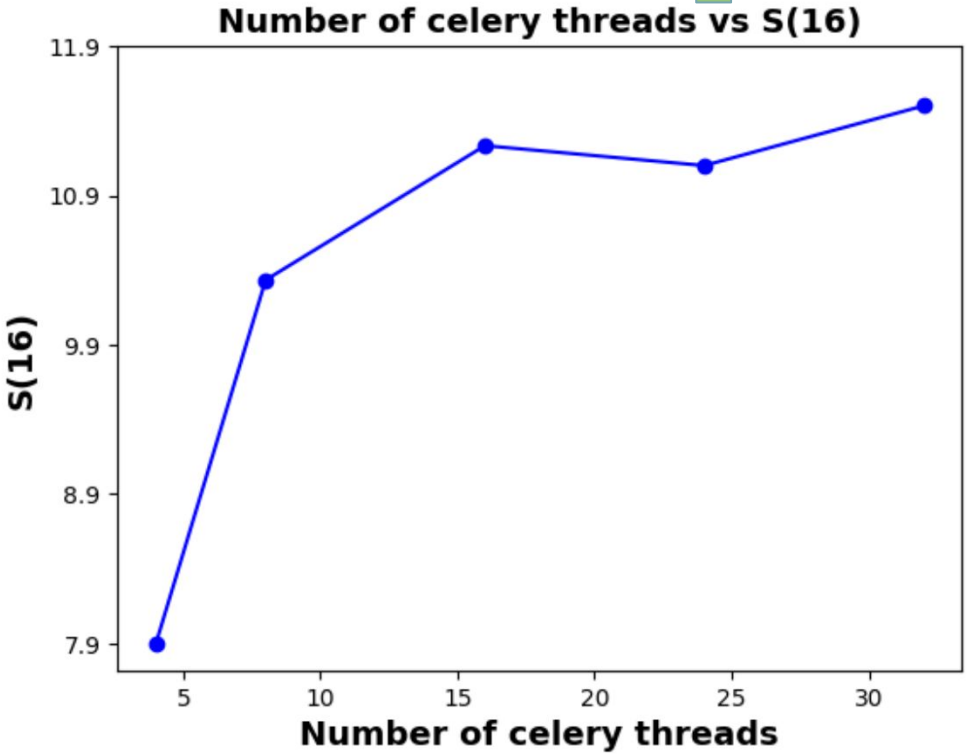

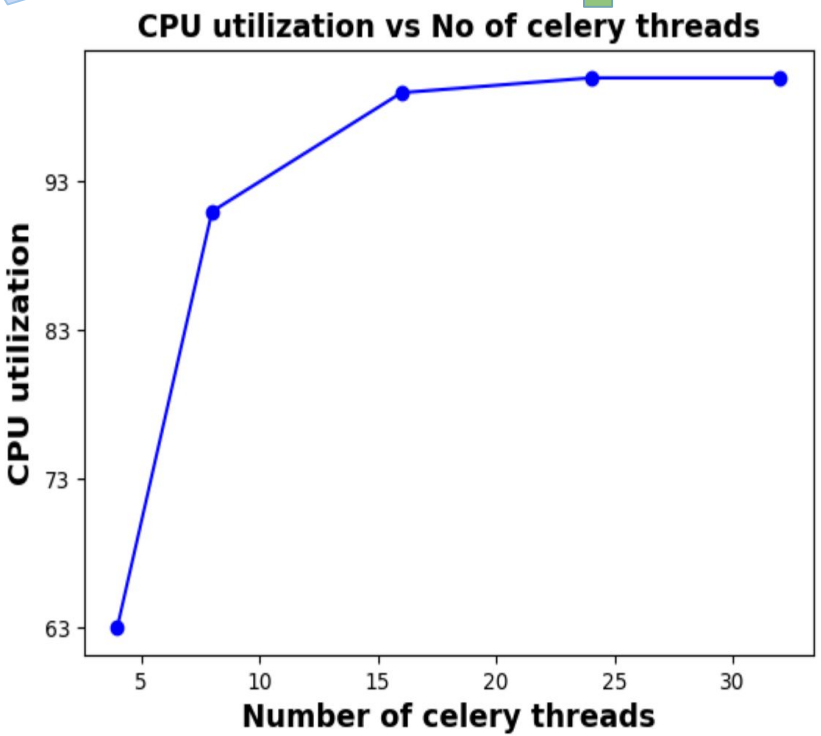

Setting CPU affinities for celery

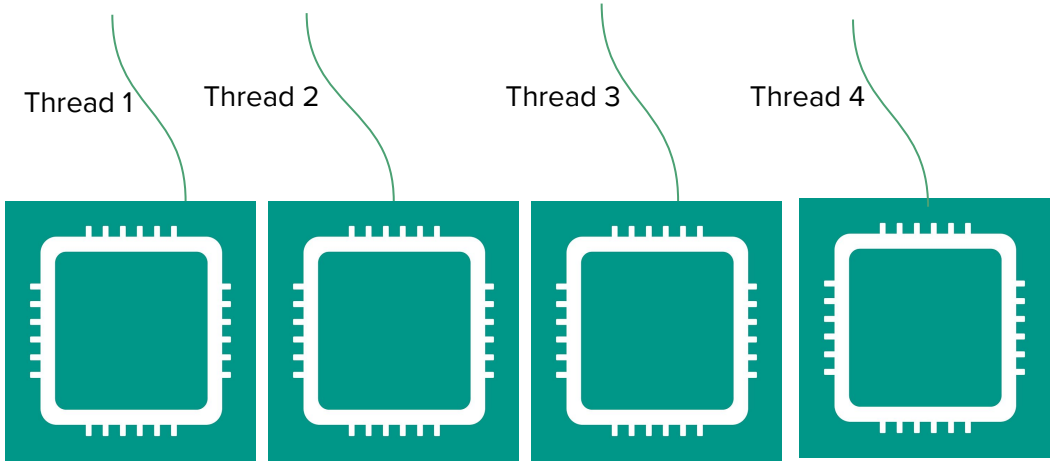

Type	Throughput (req/sec)	S(16)
Without affinity	3.37	11.23
With affinity	3.32	11

16 cores

Setting CPU affinities didn't improve throughput scalability

Increasing celery prefetch multiplier

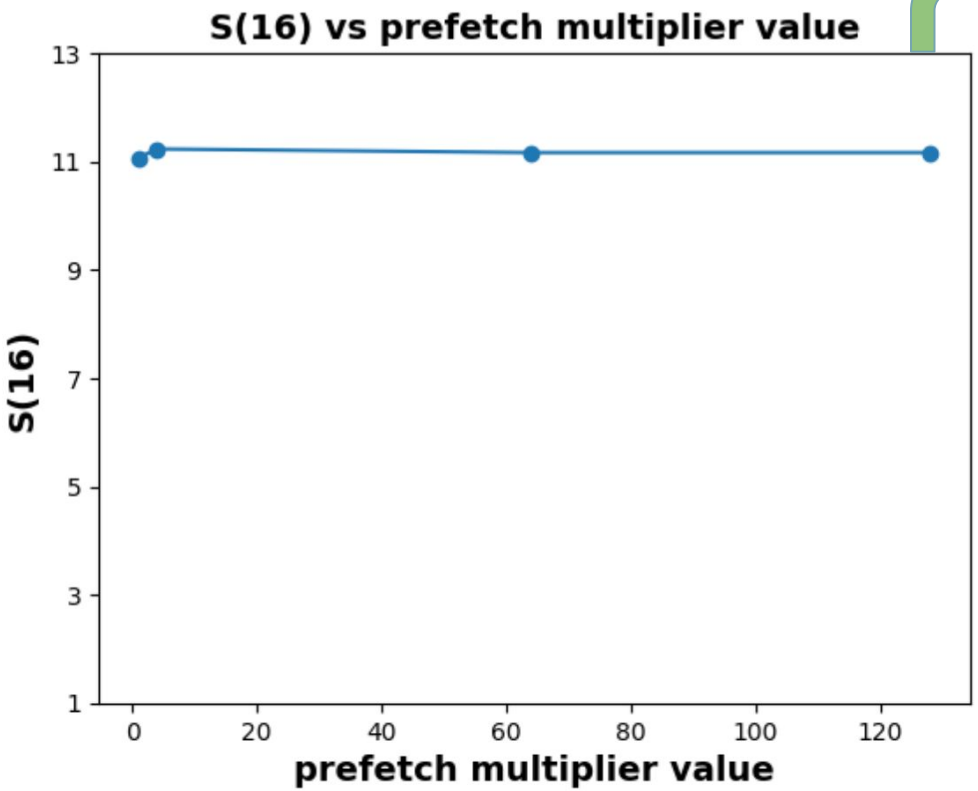

Increasing prefetch multiplier didn't improve throughput scalability

No of tasks celery picks and keeps in memory at a time

Baseline prefetch multiplier value = 4

Celery Prefetch multiplier is not bottleneck

Increasing celery broker pool limit

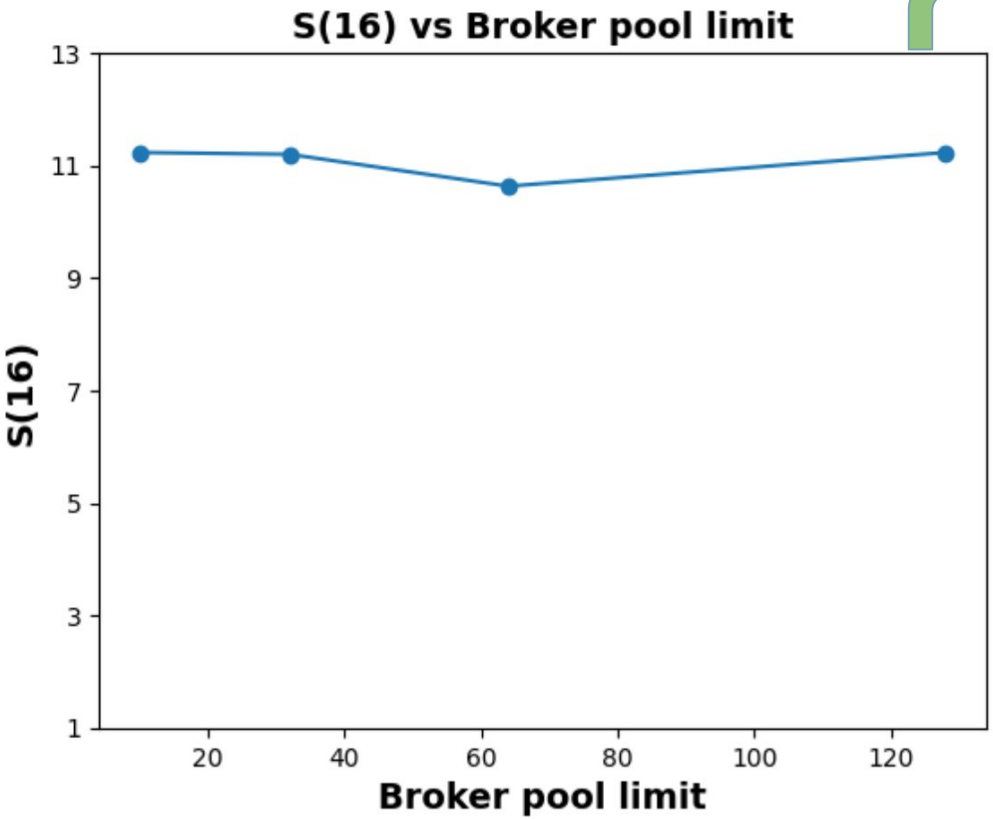

Increasing broker pool limit didn't improve throughput scalability

No of connections kept open for the broker

Baseline broker pool limit = 10

Celery Broker pool limit is not bottleneck

Using transient celery queue

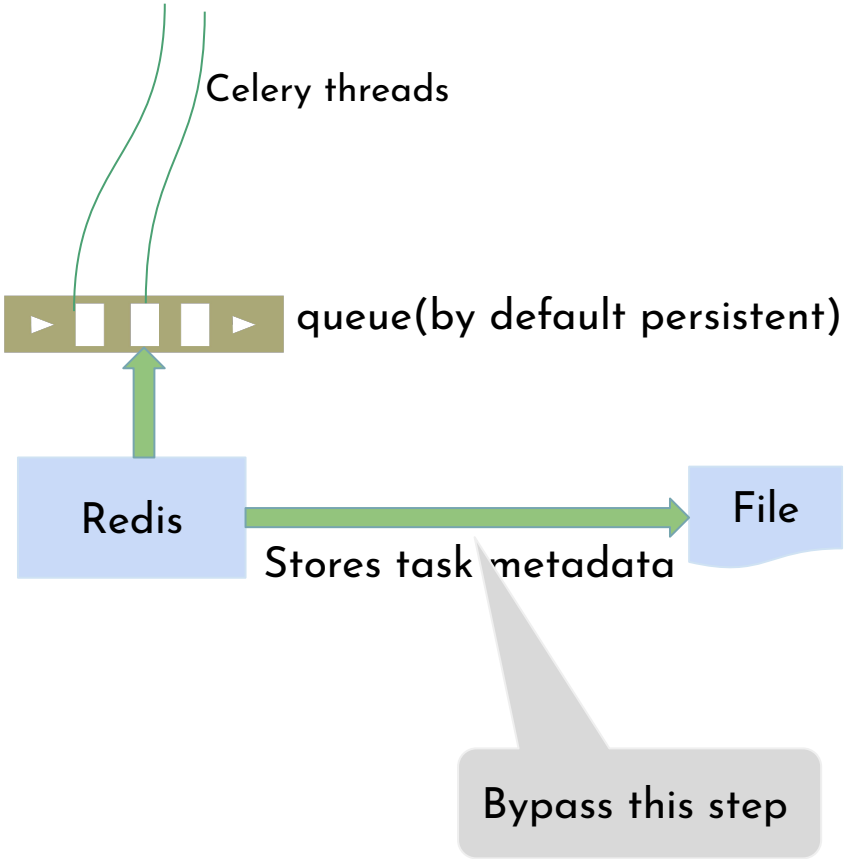

Type	Throughput (req/sec)	S(16)
Persistent	3.37	11.23
Transient	3.3	11

Making celery queue transient didn't improve throughput scalability

Disabling writing to log files

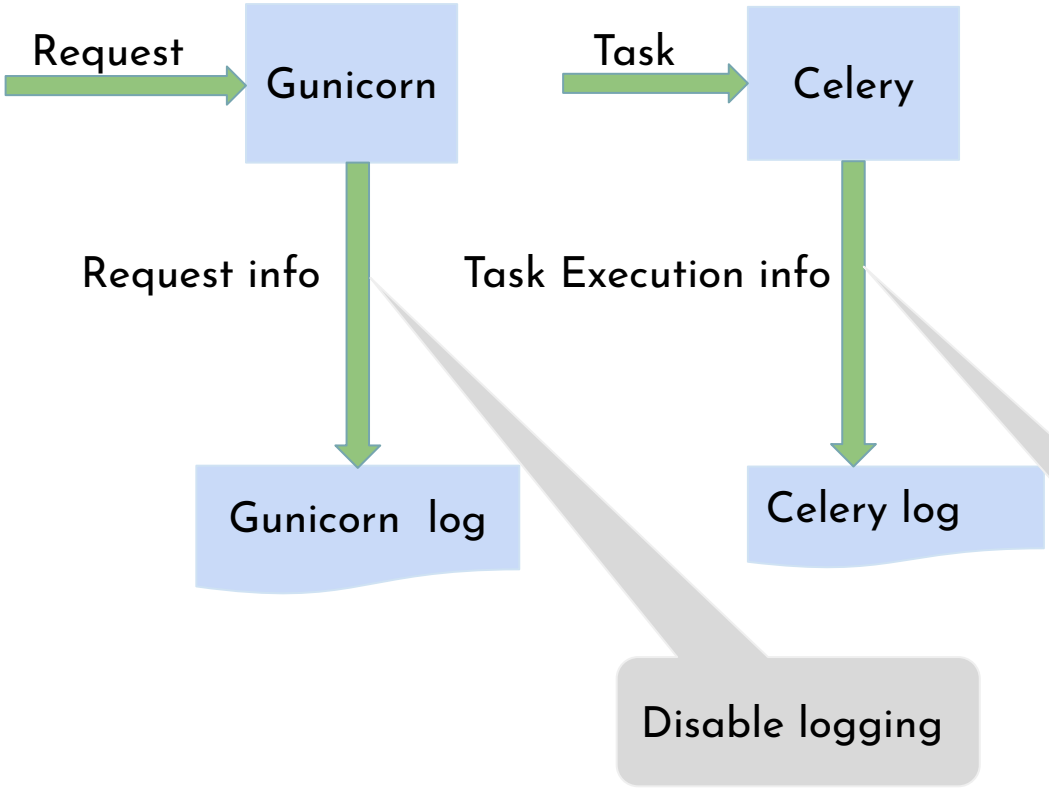

Type	Throughput (req/sec)	S(16)
Persistent	3.37	11.23
Transient	3.3	11

Disabling logs didn't improve throughput scalability